

GREEDY ALGORITHMS

INTRODUCTION

- Change Making Problem
- You Purchase a Idli & Half Tea at a canteen.
- Total bill (including GST) is 23rs.
- You pay 100rs note to the cashier.
- The balance amount is 77rs.
- What will be the approach of the cashier to pay you the balance?

INTRODUCTION

- Cashier will first pick 50rs note.
- Then he will pick 20rs note
- OR
- He will pick Two 10rs notes.
- Then he will pick a 5rs coin.
- Then finally a 2rs coin or Two 1rs coins.
- OR
- A 2rs Chocolate

INTRODUCTION

- Change Making Problem
- At least subconsciously, by millions of cashiers all over the world: give change for a specific amount n with the least number of coins of the denominations $d_1 > d_2 > \dots > d_m$ used in that locale.

INTRODUCTION

- A greedy algorithm is an **algorithmic paradigm** that follows the problem solving heuristic of making the **locally optimal choice** at each stage with the hope of finding a **global optimum**.

INTRODUCTION

- The greedy approach suggests constructing a solution through a sequence of steps, each expanding a partially constructed solution obtained so far, until a complete solution to the problem is reached.
- On each step the choice made must be:
 - **Feasible**, i.e., it has to satisfy the problem's constraints
 - **Locally Optimal**, i.e., it has to be the best local choice among all feasible choices available on that step

INTERVAL SCHEDULING PROBLEM

- We have a set of requests $\{1, 2, \dots, n\}$.
- The i th request corresponds to an interval of time starting at $s(i)$ and finishing at $f(i)$.
- We'll say that a subset of the requests is compatible if no two of them overlap in time.
- **Goal** is to accept as large a compatible subset as possible.
- Compatible sets of maximum size will


INTERVAL SCHEDULING PROBLEM

- **Designing a Greedy Algorithm using natural approaches:**
 - The basic idea in a greedy algorithm for interval scheduling is to use a simple rule to select a first request i_1 .
 - Once a request i_1 is accepted, we reject all requests that are not compatible with i_1 .
 - We then select the next request i_2 to be accepted, and again reject all requests that are not compatible with i_2 and so on.

INTERVAL SCHEDULING PROBLEM

- **Natural Rules that can be used to solve the interval scheduling problem are :**
 - Select the available request that starts earliest - that is, the one with minimal start time $s(i)$ or earliest start time.
 - This way our resource starts being used as quickly as possible.
 - If the earliest request i is for a very long interval, then by accepting request i we may have to reject a lot of requests for shorter time intervals.


INTERVAL SCHEDULING PROBLEM



INTERVAL SCHEDULING PROBLEM

- **Natural Rules that can be used to solve the interval scheduling problem are :**
 - By accepting the request that requires the smallest interval of time—namely, the request for which $f(i) - s(i)$ is as small as possible.
 - As it turns out, this is a somewhat better rule than the previous one, but it still can produce a suboptimal schedule.


INTERVAL SCHEDULING PROBLEM



INTERVAL SCHEDULING PROBLEM

- **Natural Rules that can be used to solve the interval scheduling problem are :**
 - A greedy algorithm that is based on this idea:
 - For each request, we count the number of other requests that are not compatible, and accept the request that has the fewest number of non-compatible requests.
 - In other words, we select the interval with the fewest “conflicts”.

INTERVAL SCHEDULING PROBLEM



INTERVAL SCHEDULING PROBLEM

- **Natural Rules that can be used to solve the interval scheduling problem are :**
 - A greedy rule that does lead to the optimal solution is based on a fourth idea:
 - We should accept first the request that finishes first, that is, the request i for which $f(i)$ is as small as possible.

INTERVAL SCHEDULING PROBLEM


- **Interval Scheduling Algorithm**

```
Initially let  $R$  be the set of all requests, and let  $A$  be empty
  While  $R$  is not yet empty
    Choose a request  $i \in R$  that has the smallest finishing
    time
    Add request  $i$  to  $A$ 
    Delete all requests from  $R$  that are not compatible with
    request  $i$ 
  EndWhile
Return the set  $A$  as the set of accepted requests
```


INTERVAL SCHEDULING PROBLEM

- Interval Scheduling Example with Smallest Finish Time (Greedy)**


Intervals numbered in order



INTERVAL SCHEDULING PROBLEM

- Interval Scheduling Example with Smallest Finish Time (Greedy)

Selecting interval 1




$i=1$

$j=2,6$

$A=\{1\}$


INTERVAL SCHEDULING PROBLEM

- Interval Scheduling Example with Smallest Finish Time (Greedy)




INTERVAL SCHEDULING PROBLEM

- Interval Scheduling Example with Smallest Finish Time (Greedy)



INTERVAL SCHEDULING PROBLEM

- Interval Scheduling Example with Smallest Finish Time (Greedy)



INTERVAL SCHEDULING PROBLEM

- Implementation and Running Time of Interval Scheduling:
 - The running time of the algorithm shown is $O(n \log n)$ as follows.
 - We begin by sorting the n requests in order of finishing time and labeling them in this order; that is, we will assume that $f(i) \leq f(j)$ when $i < j$.
 - This takes time $O(n \log n)$.
 - In an additional $O(n)$ time, we construct an array $S[1 \dots n]$ with the property that $S[i]$ contains the value $s(i)$.
 - We now select requests by processing the intervals in order of increasing $f(i)$.

INTERVAL SCHEDULING PROBLEM

- We always select the first interval; we then iterate through the intervals in order until reaching the first interval j for which $s(j) \geq f(1)$; we then select this one as well.
- More generally, if the most recent interval we've selected ends at time f , we continue iterating through subsequent intervals until we reach the first j for which $s(j) \geq f$.
- In this way, we implement the greedy algorithm analyzed above in one pass through the intervals, spending constant time per interval. Thus this part of the algorithm takes time $O(n)$.

SCHEDULING TO MINIMIZE LATENESS

- Consider again a situation in which we have a single resource and a set of n requests to use the resource for an interval of time.
- Assume that the resource is available starting at time S and each request is now more flexible.
- Instead of a start time and finish time, the request i has a deadline d_i , and it requires a contiguous time interval of length t_i , but it is willing to be scheduled at any time before the

SCHEDULING TO MINIMIZE LATENESS

- Suppose that we plan to satisfy each request, but we are allowed to let certain requests run late.
- Thus, beginning at our overall start time S , we will assign each request i an interval of time of length t_i ;
 - Let us denote this interval by $[s(i), f(i)]$, with $f(i)$ as
 - $f(i) = s(i) + t_i$.
 - where, $s(i)$ = start time of the interval i
 - $f(i)$ = finish time of the interval i
 - t_i = time length of the interval i

SCHEDULING TO MINIMIZE LATENESS

- The algorithm must determine a start time (and hence a finish time) for each interval.
- We say that a request i is late if it misses the deadline, that is, if $f(i) > d_i$.
- The lateness of such a request i is defined to be
 - $l_i = f(i) - d_i$
- We will say that $l_i = 0$ if request i is not late.
- The goal in our new optimization problem will be to schedule all requests, using non overlapping intervals, so as to

SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

SCHEDULING TO MINIMIZE LATENESS


	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

$d_3 = 9$

0 1 2 3 4 5 6 7 8 9 10


SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15




SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15




SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



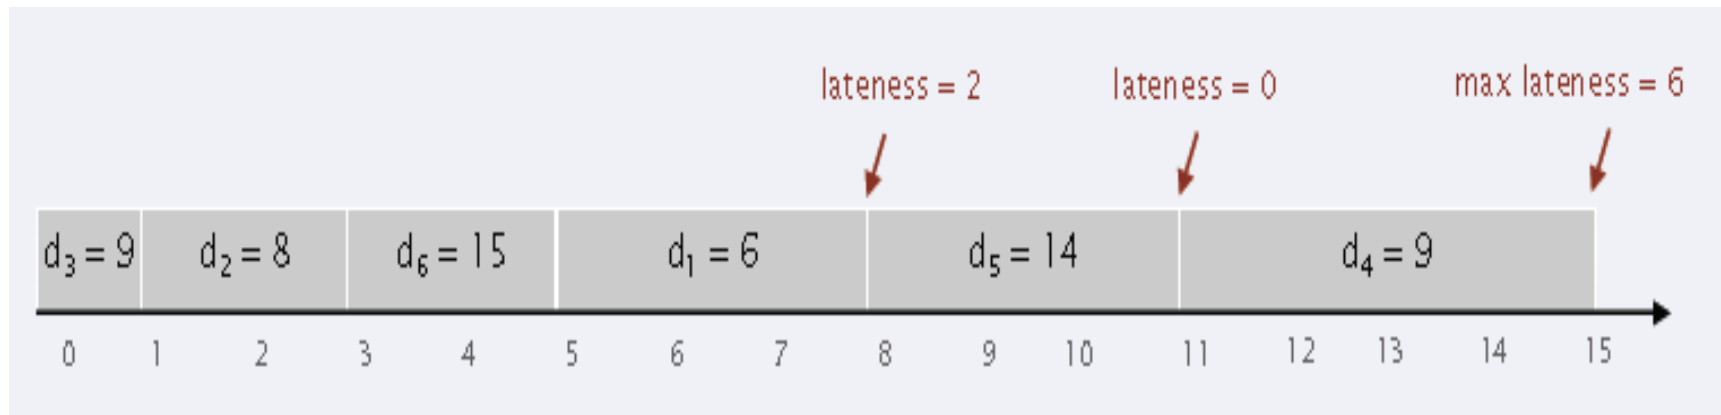
SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



SCHEDULING TO MINIMIZE LATENESS

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15




SCHEDULING TO MINIMIZE LATENESS

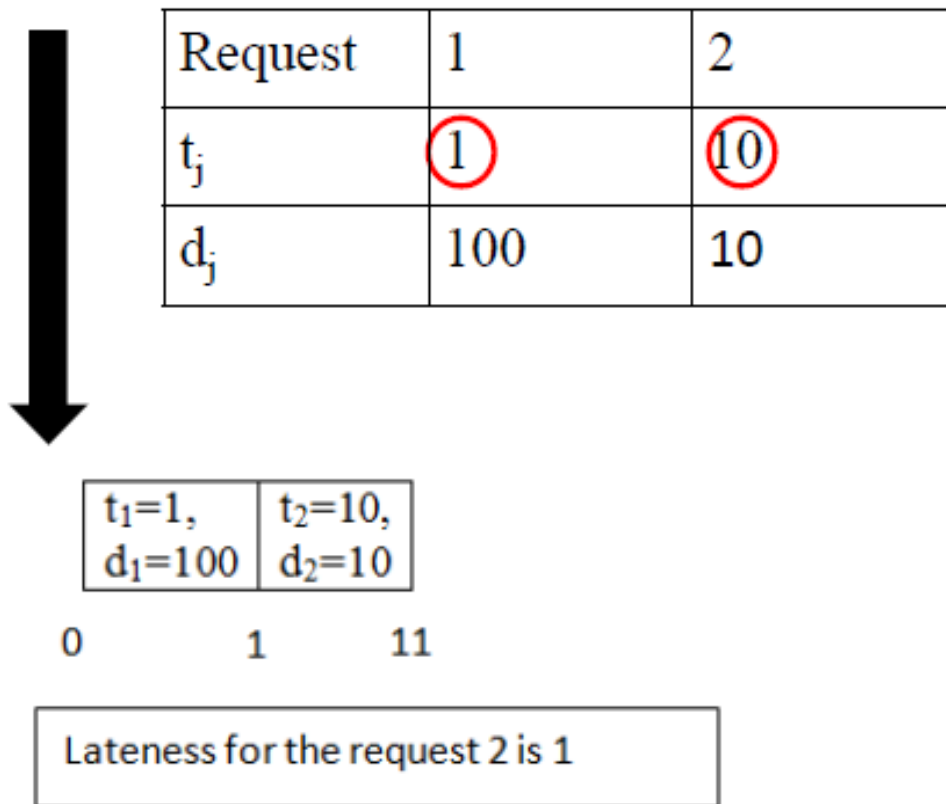
- Natural Approaches to the Problem.
 - Schedule the jobs in order of

incr

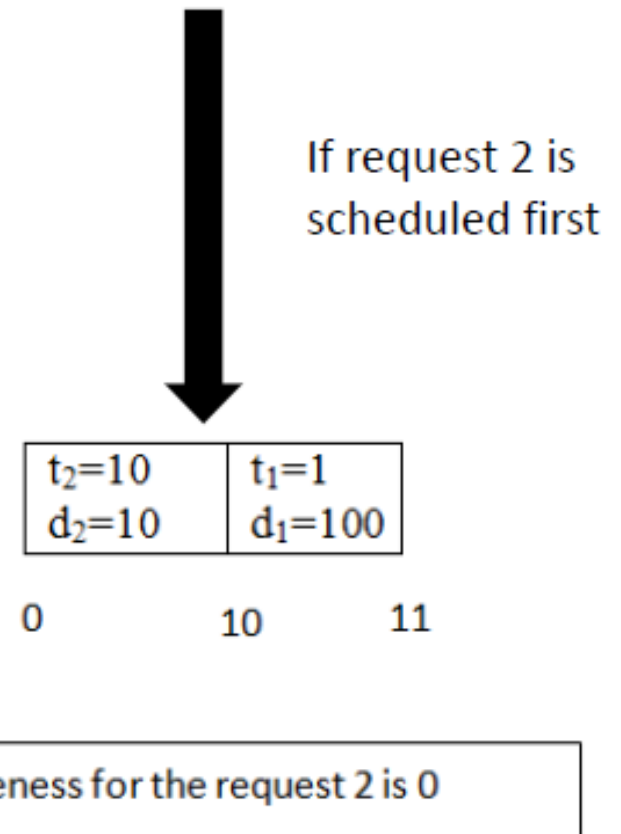
Request	1	2
t_j	1	3
d_j	2	5



SCHEDULING TO MINIMIZE LATENESS



Counter example



SCHEDULING TO MINIMIZE LATENESS

- Natural Approaches to the Problem.
 - Slack Time: $d_i - t_i$ is very small.

Request	1	2
t_j	1	3
d_j	2	5
Sl_i (Slack)	1	2

$sl_1=1$	$sl_2=2$
$d_1=2$	$d_2=5$

0 1 4

SCHEDULING TO MINIMIZE LATENESS

Counter example



Request	1	2
t_j	1	10
d_j	2	10
sl_i	1	0



If request 1
Was scheduled
first

$sl_2=0$ $d_2=10$	$sl_1=1$ $d_1=2$
----------------------	---------------------

0 10 11

Lateness for the request 1 is 9

$sl_1=1$ $d_1=2$	$sl_2=0$ $d_2=10$
---------------------	----------------------

0 1 11

Lateness for the request 2 is 1

SCHEDULING TO MINIMIZE LATENESS – GOAL REVISITED

- The goal in our new optimization problem will be to schedule all requests, using non-overlapping intervals, so as to minimize the maximum lateness,
 $L = \max_i l_i$

SCHEDULING TO MINIMIZE LATENESS – GOAL REVISITED

Algorithm

Order the jobs in order of their deadlines

Assume for simplicity of notation that $d_1 \leq \dots \leq d_n$

Initially, $f = s$

Consider the jobs $i = 1, \dots, n$ in this order

 Assign job i to the time interval from $s(i) = f$ to $f(i) = f + t_i$

 Let $f = f + t_i$


End

Return the set of scheduled intervals $[s(i), f(i)]$ for $i = 1, \dots, n$

SCHEDULING TO MINIMIZE LATENESS – EXAMPLE

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

1. $i=1$ Add Job J1 to the schedule.



$$s(1) = f = 0$$

$$f(1) = s(1) + t(1) = 0 + 3 = 3$$

$$f = f + t(1) = 0 + 3 = 3$$

SCHEDULING TO MINIMIZE LATENESS – EXAMPLE

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

2. $i=2$ Add Job J2 to the schedule.

J1 $d_1=6$	J2 $d_2=8$
------------	------------

0 3 5

$$s(2) = f = 3$$

$$f(2) = s(2) + t(2) = 3 + 2 = 5$$

$$f = f + t(2) = 3 + 2 = 5$$

SCHEDULING TO MINIMIZE LATENESS – EXAMPLE

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

3. $i=3$ Add Job J3 to the schedule.

J1 d1=6	J2 d2=8	J3 d3=9
0	3	5

$$s(3) = f = 5$$

$$f(3) = s(3) + t(3) = 5 + 1 = 6$$

$$f = f + t(3) = 5 + 1 = 6$$

SCHEDULING TO MINIMIZE LATENESS – EXAMPLE


	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

4. $i=4$ Add Job J4 to the schedule.

$$s(4) = f = 6$$

$$f(4) = s(4) + t(4) = 6 + 4 = 10$$

$$f = f + t(4) = 6 + 4 = 10$$



SCHEDULING TO MINIMIZE LATENESS – EXAMPLE

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

5. $i=5$ Add Job J5 to the schedule.

$$\begin{aligned} s(5) &= f = 10 \\ f(5) &= s(5) + t(5) = 10 + 3 = 13 \\ f &= f + t(5) = 10 + 3 = 13 \end{aligned}$$


J1 d1=6	J2 d2=8	J3 d3=9	J4 d4=9	J5 d5=14	
0	3	5	6	10	13

SCHEDULING TO MINIMIZE LATENESS – EXAMPLE

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15

6. $i=6$ Add Job J6 to the schedule.

$$\begin{aligned} s(6) &= f = 13 \\ f(6) &= s(6) + t(6) = 13 + 2 = 15 \\ f &= f + t(6) = 13 + 2 = 15 \end{aligned}$$



OPTIMAL CACHING-PROBLEM


- Caching.
 - Cache with capacity to store k items.
 - Sequence of m item requests d_1, d_2, \dots, d_m .
 - Cache hit: item already in cache when requested.
 - Cache miss: item not already in cache when requested: must bring requested item into cache, and evict some existing item, if full

OPTIMAL CACHING-PROBLEM

- Goal - Eviction schedule that minimizes number of cache misses.
 - Example
 - $k = 2$, initial cache = a b, requests: a, b, c, b, c, a, a, b.
 - Optimal eviction schedule: 2 cache misses.

OPTIMAL CACHING-PROBLEM

red = cache miss



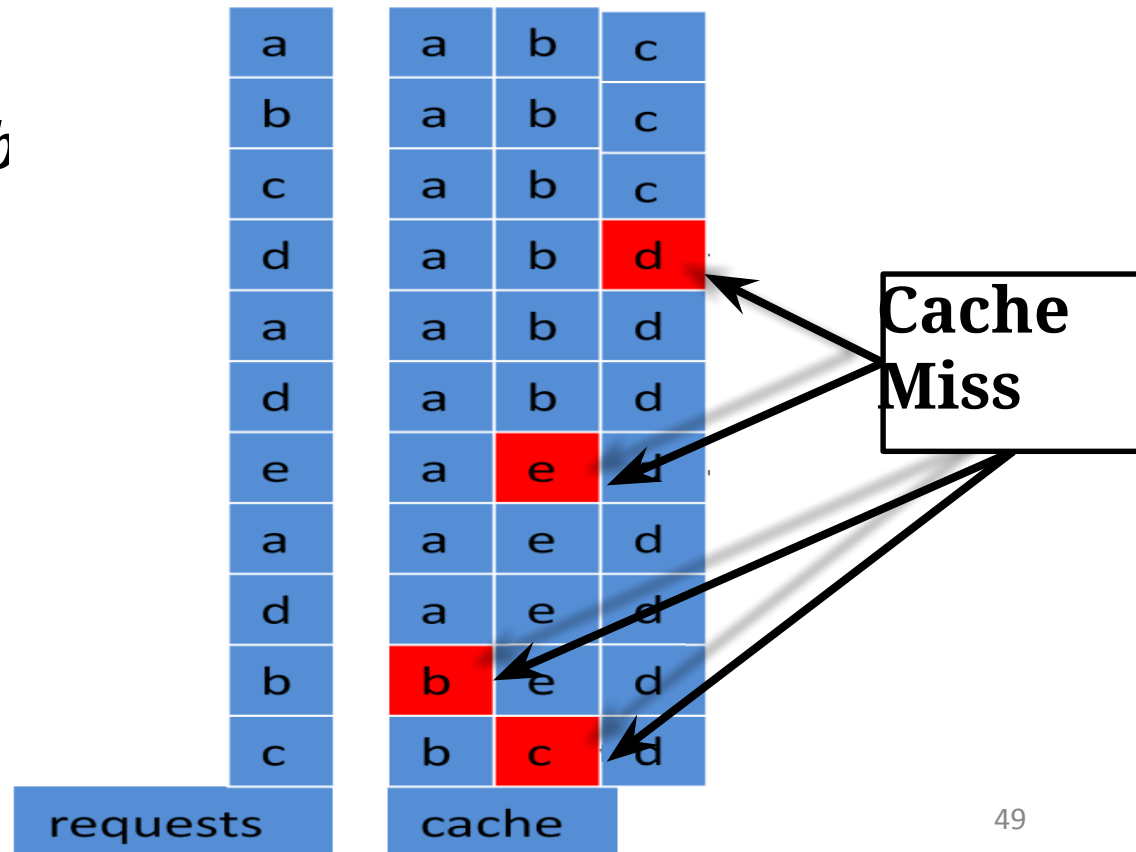
OPTIMAL CACHING- ALGORITHM -Farthest-in- Future

Algorithm

When d_i needs to be brought into the cache,
evict the item that is needed the farthest into the future

Example

$a, b, c, d, a, d, e, a, d, b$
 $k=3$



THANK YOU