

**M.S. Ramaiah Institute of Technology**  
**(Autonomous Institute, Affiliated to VTU)**  
**Department of Computer Science and Engineering**

**Course Name: Distributed Systems**

**Course Code: CSE751**

**Credits: 3:0:0**

**Term: September – December 2020**

---

Faculty:  
Sini Anna Alex

# Cuts of a Distributed Computation

---

“In the space-time diagram of a distributed computation, a *cut* is a zigzag line joining one arbitrary point on each process line.”

- A cut slices the space-time diagram, and thus the set of events in the distributed computation, into a PAST and a FUTURE.
- The PAST contains all the events to the left of the cut and the FUTURE contains all the events to the right of the cut.
- For a cut  $C$ , let  $PAST(C)$  and  $FUTURE(C)$  denote the set of events in the PAST and FUTURE of  $C$ , respectively.
- Every cut corresponds to a global state and every global state can be graphically represented as a cut in the computation's space-time diagram.
- Cuts in a space-time diagram provide a powerful graphical aid in representing and reasoning about global states of a computation.

# Cuts of a Distributed Computation





Illustration of cuts in a distributed execution

# Cuts of a Distributed Computation

---

- In a consistent cut, every message received in the PAST of the cut was sent in the PAST of that cut. (In Figure 2.3, cut  $C_2$  is a consistent cut.)
- All messages that cross the cut from the PAST to the FUTURE are in transit in the corresponding consistent global state.
- A cut is *inconsistent* if a message crosses the cut from the FUTURE to the PAST. (In Figure 2.3, cut  $C_1$  is an inconsistent cut.)

# Illustration of consistent states





# Past and Future Cones of an Event

---

## Past Cone of an Event

- An event  $e_j$  could have been affected only by all events  $e_i$  such that  $e_i \rightarrow e_j$ .
- In this situation, all the information available at  $e_i$  could be made accessible at  $e_j$ .
- All such events  $e_i$  belong to the past of  $e_j$ .

Let  $Past(e_j)$  denote all events in the past of  $e_j$  in a computation  $(H, \rightarrow)$ . Then,

$$Past(e_j) = \{e_i | \forall e_i \in H, e_i \rightarrow e_j\}.$$

- Figure 2.4 (next slide) shows the past of an event  $e_j$ .

# Jard–Jourdan’s adaptive technique

---

Direct-dependency technique(The Fowler–Zwaenepoel) does not allow the transitive dependencies to be captured in real time during the execution of processes.

In addition, a **process must observe an event after receiving a message but before sending out any message**. Otherwise, during the reconstruction of a **vector timestamp from the direct-dependency vectors, all the causal dependencies will not be captured**.

If events occur very frequently, this technique will require recording the history of a large number of events.

# Jard–Jourdan's adaptive technique

---

In the Jard–Jourdan's technique, events can be adaptively observed while maintaining the capability of retrieving all the causal dependencies of an observed event. (Observing an event means recording of the information about its dependencies.)

This method uses the idea that when an observed event  $e$  records its dependencies, then events that follow can determine their transitive dependencies of an observed event.



# Jard–Jourdan’s adaptive technique

---

Jard–Jourdan defined a *pseudo-direct* relation  $\ll$  on the events of a distributed computation as follows:

If events  $e_i$  and  $e_j$  happen at process  $p_i$  and  $p_j$ , respectively, then  $e_j \ll e_i$  iff there exists a path of message transfers that starts after  $e_j$  on the process  $p_j$  and ends before  $e_i$  on the process  $p_i$  such that there is no observed event on the path.

# Technique is implemented using the following mechanism

---

Initially, at a process  $p_i$ :  $p\_vt_i = \{(i, 0)\}$ .

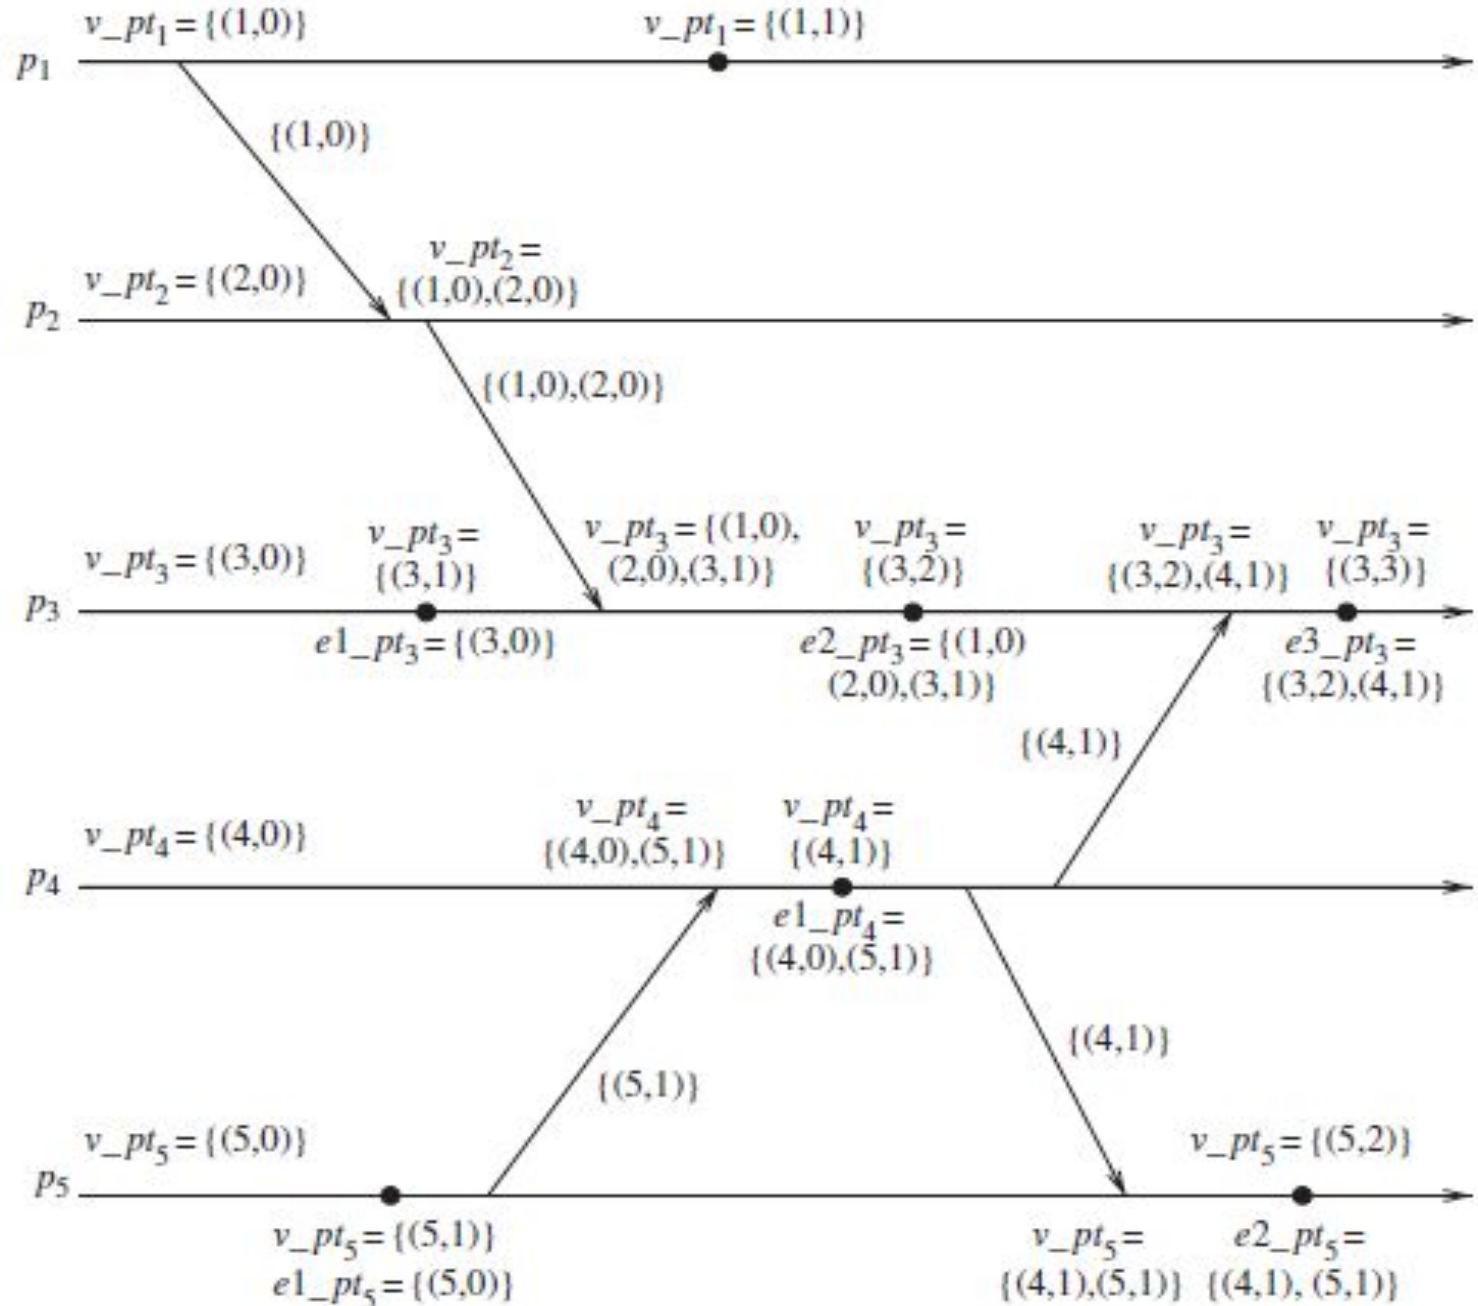
Let  $p\_vt_i = \{(i_1, v_1), \dots, (i, v), \dots, (i_n, v_n)\}$  denote the current partial vector clock at process  $p_i$ . Let  $e\_vt_i$  be a variable that holds the timestamp of the observed event.

Whenever an event is observed at process  $p_i$ , the contents of the partial vector clock  $p\_vt_i$  are transferred to  $e\_vt_i$  and  $p\_vt_i$  is reset and updated as follows:


$$\begin{aligned} e\_vt_i &= \{(i_1, v_1), \dots, (i, v), \dots, (i_n, v_n)\} \\ p\_vt_i &= \{(i, v + 1)\}. \end{aligned}$$

When process  $p_j$  sends a message to  $p_i$ , it piggybacks the current value of  $p\_vt_j$  in the message.


## Vector clocks progress in the Jard–Jourdan technique



# Jard–Jourdan's adaptive technique



# Identify the consistent and inconsistent cuts in distributed computation



# Text Book

---

Ajay D. Kshemkalyani and MukeshSinghal, *Distributed Computing: Principles, Algorithms, and Systems*, Cambridge University Press.



Thank you