

**M.S. Ramaiah Institute of Technology**  
**(Autonomous Institute, Affiliated to VTU)**  
**Department of Computer Science and Engineering**

**Course Name: Distributed Systems**

**Course Code: CSE20/CSE751**

**Credits: 3:0:0**

**Term: Oct 2021-Feb 2022**

---

Faculty:  
Sini Anna Alex

L1:  $S_i$  has received a message with timestamp larger than  $(ts_i, i)$  from all other sites.

L2:  $S_i$ 's request is at the top of  $request\_queue_i$ .

# Correctness

---

**Theorem: Lamport's algorithm achieves mutual exclusion.**

**Proof:**

- Proof is by contradiction. Suppose two sites  $S_i$  and  $S_j$  are executing the CS concurrently. For this to happen conditions L1 and L2 must hold at both the sites *concurrently*.
- This implies that at some instant in time, say  $t$ , both  $S_i$  and  $S_j$  have their own requests at the top of their *request\_queues* and condition L1 holds at them. Without loss of generality, assume that  $S_i$ 's request has smaller timestamp than the request of  $S_j$ .
- From condition L1 and FIFO property of the communication channels, it is clear that at instant  $t$  the request of  $S_i$  must be present in *request\_queue\_j* when  $S_j$  was executing its CS. This implies that  $S_j$ 's own request is at the top of its own *request\_queue* when a smaller timestamp request,  $S_i$ 's request, is present in the *request\_queue\_j* – a contradiction!

# Correctness

---

**Theorem:** Lamport's algorithm is fair.

**Proof:**

- The proof is by contradiction. Suppose a site  $S_i$ 's request has a smaller timestamp than the request of another site  $S_j$  and  $S_j$  is able to execute the CS before  $S_i$ .
- For  $S_j$  to execute the CS, it has to satisfy the conditions L1 and L2. This implies that at some instant in time say  $t$ ,  $S_j$  has its own request at the top of its queue and it has also received a message with timestamp larger than the timestamp of its request from all other sites.
- But *request\_queue* at a site is ordered by timestamp, and according to our assumption  $S_i$  has lower timestamp. So  $S_i$ 's request must be placed ahead of the  $S_j$ 's request in the *request\_queue<sub>j</sub>*. This is a contradiction!



# An optimization

---

- In Lamport's algorithm, REPLY messages can be omitted in certain situations. For example, if site  $S_j$  receives a REQUEST message from site  $S_i$  after it has sent its own REQUEST message with timestamp higher than the timestamp of site  $S_i$ 's request, then site  $S_j$  need not send a REPLY message to site  $S_i$ .
- This is because when site  $S_i$  receives site  $S_j$ 's request with timestamp higher than its own, it can conclude that site  $S_j$  does not have any smaller timestamp request which is still pending.
- With this optimization, Lamport's algorithm requires between  $3(N - 1)$  and  $2(N - 1)$  messages per CS execution.

# Ricart-Agrawala Algorithm

---

- The Ricart-Agrawala algorithm assumes the communication channels are FIFO. The algorithm uses two types of messages: REQUEST and REPLY.
- A process sends a REQUEST message to all other processes to request their permission to enter the critical section. A process sends a REPLY message to a process to give its permission to that process.
- Processes use Lamport-style logical clocks to assign a timestamp to critical section requests and timestamps are used to decide the priority of requests.
- Each process  $p_i$  maintains the Request-Deferred array,  $RD_i$ , the size of which is the same as the number of processes in the system.
- Initially,  $\forall i \forall j: RD_i[j]=0$ . Whenever  $p_i$  defer the request sent by  $p_j$ , it sets  $RD_i[j]=1$  and after it has sent a REPLY message to  $p_j$ , it sets  $RD_i[j]=0$ .



# Description of the Algorithm

---

## Requesting the critical section:

- (a) When a site  $S_i$  wants to enter the CS, it broadcasts a timestamped REQUEST message to all other sites.
- (b) When site  $S_j$  receives a REQUEST message from site  $S_i$ , it sends a REPLY message to site  $S_i$  if site  $S_j$  is neither requesting nor executing the CS, or if the site  $S_j$  is requesting and  $S_i$ 's request's timestamp is smaller than site  $S_j$ 's own request's timestamp. Otherwise, the reply is deferred and  $S_j$  sets  $RD_j[i]=1$

## Executing the critical section:

- (c) Site  $S_i$  enters the CS after it has received a REPLY message from every site it sent a REQUEST message to.

# Description of the Algorithm

---


## Releasing the critical section:

- (d) When site  $S_i$  exits the CS, it sends all the deferred REPLY messages:  $\forall j$  if  $RD_i[j]=1$ , then send a REPLY message to  $S_j$  and set  $RD_i[j]=0$ .

## Notes:

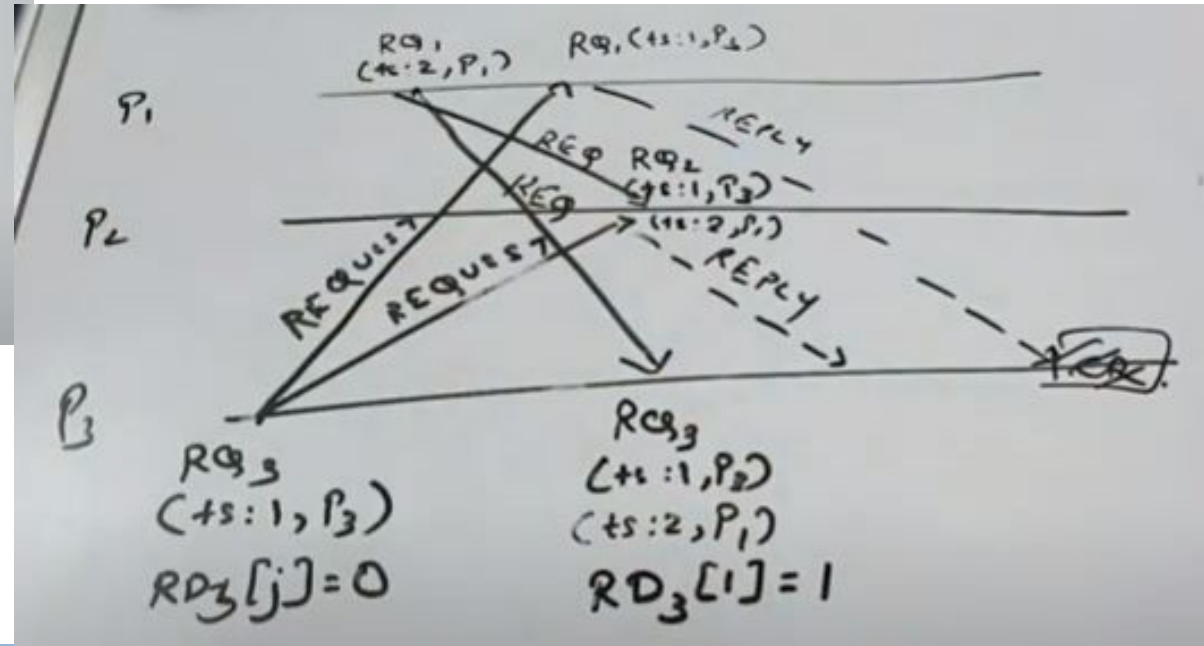

- When a site receives a message, it updates its clock using the timestamp in the message.
- When a site takes up a request for the CS for processing, it updates its local clock and assigns a timestamp to the request.

# Example






# Example



# Example



Thank you