

# Chapter 14: Consensus and Agreement

Ajay Kshemkalyani and Mukesh Singhal


Distributed Computing: Principles, Algorithms, and Systems

Cambridge University Press

# Assumptions

## System assumptions

- Failure models
- Synchronous/ Asynchronous communication
- Network connectivity
- Sender identification
- Channel reliability
- Authenticated vs. non-authenticated messages
- Agreement variable



# Problem Specifications

Byzantine Agreement (single source has an initial value)

**Agreement:** All non-faulty processes must agree on the same value.

**Validity:** If the source process is non-faulty, then the agreed upon value by all the non-faulty processes must be the same as the initial value of the source.

**Termination:** Each non-faulty process must eventually decide on a value.

Consensus Problem (all processes have an initial value)

**Agreement:** All non-faulty processes must agree on the same (single) value.

**Validity:** If all the non-faulty processes have the same initial value, then the agreed upon value by all the non-faulty processes must be that same value.

**Termination:** Each non-faulty process must eventually decide on a value.

Interactive Consistency (all processes have an initial value)

**Agreement:** All non-faulty processes must agree on the same array of values  $A[v_1 \dots v_n]$ .

**Validity:** If process  $i$  is non-faulty and its initial value is  $v_i$ , then all non-faulty processes agree on  $v_i$  as the  $i$ th element of the array  $A$ . If process  $j$  is faulty, then the non-faulty processes can agree on any value for  $A[j]$ .

**Termination:** Each non-faulty process must eventually decide on the array  $A$ .

These problems are equivalent to one another! Show using reductions.

# Overview of Results

Failure mode	Synchronous system (message-passing and shared memory)	Asynchronous system (message-passing and shared memory)
No failure	agreement attainable; common knowledge also attainable	agreement attainable; concurrent common knowledge attainable
Crash failure	agreement attainable $f < n$ processes $\Omega(f + 1)$ rounds	agreement not attainable
Byzantine failure	agreement attainable $f \leq \lfloor (n - 1)/3 \rfloor$ Byzantine processes $\Omega(f + 1)$ rounds	agreement not attainable

**Table:** Overview of results on agreement.  $f$  denotes number of failure-prone processes.  $n$  is the total number of processes.

In a failure-free system, consensus can be attained in a straightforward manner

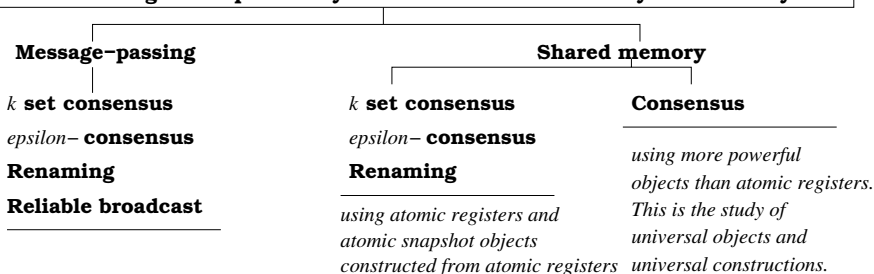
# Some Solvable Variants of the Consensus Problem in Async Systems

Solvable Variants	Failure model and overhead	Definition
Reliable broadcast	crash failures, $n > f$ (MP)	Validity, Agreement, Integrity conditions
$k$ -set consensus	crash failures. $f < k < n$ . (MP and SM)	size of the set of values agreed upon must be less than $k$
$\epsilon$ -agreement	crash failures $n \geq 5f + 1$ (MP)	values agreed upon are within $\epsilon$ of each other
Renaming	up to $f$ fail-stop processes, $n \geq 2f + 1$ (MP) Crash failures $f \leq n - 1$ (SM)	select a unique name from a set of names

**Table:** Some solvable variants of the agreement problem in asynchronous system. The overhead bounds are for the given algorithms, and not necessarily tight bounds for the problem.

# Solvable Variants of the Consensus Problem in Async Systems

## Circumventing the impossibility results for consensus in asynchronous systems



# Consensus Algorithm for Crash Failures (MP, synchronous)

- Up to  $f$  ( $< n$ ) crash failures possible.
- In  $f + 1$  rounds, at least one round has no failures.
- Now justify: agreement, validity, termination conditions are satisfied.
- Complexity:  $O(f + 1)n^2$  messages
- $f + 1$  is lower bound on number of rounds

(global constants)

integer:  $f$ ; // maximum number of crash failures tolerated

(local variables)

integer:  $x \leftarrow$  local value;

(1) Process  $P_i$  ( $1 \leq i \leq n$ ) executes the Consensus algorithm for up to  $f$  crash failures:

(1a) **for** round **from** 1 **to**  $f + 1$  **do**

(1b)     **if** the current value of  $x$  has not been broadcast **then**

(1c)             **broadcast**( $x$ );


(1d)      $y_j \leftarrow$  value (if any) received from process  $j$  in this round;

(1e)      $x \leftarrow \min(x, y_j)$ ;

(1f) **output**  $x$  as the consensus value.

# Upper Bound on Byzantine Processes (sync)

Agreement impossible when  $f = 1, n = 3$ .

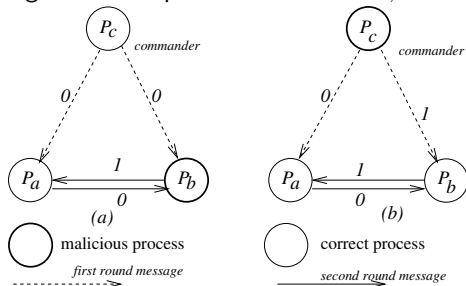


- Taking simple majority decision does not help because loyal commander  $P_a$  cannot distinguish between the possible scenarios (a) and (b);
- hence does not know which action to take.
- Proof using induction that problem solvable if  $f \leq \lfloor \frac{n-1}{3} \rfloor$ . See text.



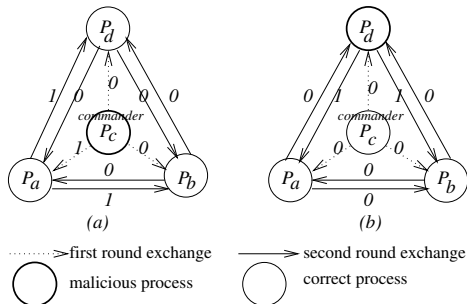
# Upper Bound on Byzantine Processes (sync)

Agreement impossible when  $f = 1, n = 3$ .



- Taking simple majority decision does not help because loyal commander  $P_a$  cannot distinguish between the possible scenarios (a) and (b);
- hence does not know which action to take.
- Proof using induction that problem solvable if  $f \leq \lfloor \frac{n-1}{3} \rfloor$ . See text.

# Consensus Solvable when $f = 1, n = 4$



- There is no ambiguity at any loyal commander, when taking majority decision
- Majority decision is over 2nd round messages, and 1st round message received directly from commander-in-chief process.

# Byzantine Generals (recursive formulation), (sync, msg-passing)

(variables)  
**boolean:**  $v \leftarrow$  initial value;  
**integer:**  $f \leftarrow$  maximum number of malicious processes,  $\leq \lfloor (n - 1)/3 \rfloor$ ;  
 (message type)  
 $Oral\_Msg(v, Dests, List, faulty)$ , where  
 $v$  is a boolean,  
 $Dests$  is a set of destination process ids to which the message is sent,  
 $List$  is a list of process ids traversed by this message, ordered from most recent to earliest,  
 $faulty$  is an integer indicating the number of malicious processes to be tolerated.

$Oral\_Msg(f)$ , where  $f > 0$ :

- ① The algorithm is initiated by the Commander, who sends his source value  $v$  to all other processes using a  $OM(v, N, \langle i \rangle, f)$  message. The commander returns his own value  $v$  and terminates.
- ② **[Recursion unfolding:]** For each message of the form  $OM(v_j, Dests, List, f')$  received in this round from some process  $j$ , the process  $i$  uses the value  $v_j$  it receives from the source, and using that value, acts as a new source. (If no value is received, a default value is assumed.)  
 To act as a new source, the process  $i$  initiates  $Oral\_Msg(f' - 1)$ , wherein it sends  
 $OM(v_j, Dests - \{i\}, concat(\langle i \rangle, L), (f' - 1))$   
 to destinations not in  $concat(\langle i \rangle, L)$   
 in the next round.
- ③ **[Recursion folding:]** For each message of the form  $OM(v_j, Dests, List, f')$  received in Step 2, each process  $i$  has computed the agreement value  $v_k$ , for each  $k$  not in  $List$  and  $k \neq i$ , corresponding to the value received from  $P_k$  after traversing the nodes in  $List$ , at one level lower in the recursion. If it receives no value in this round, it uses a default value. Process  $i$  then uses the value  $majority_{k \notin List, k \neq i}(v_j, v_k)$  as the agreement value and returns it to the next higher level in the recursive invocation.

$Oral\_Msg(0)$ :

- ① **[Recursion unfolding:]** Process acts as a source and sends its value to each other process.
- ② **[Recursion folding:]** Each process uses the value it receives from the other sources, and uses that value as the agreement value. If no value is received, a default value is assumed.

# Relationship between # Messages and Rounds

round number	a message has already visited	aims to tolerate these many failures	and each message gets sent to	total number of messages in round
1	1	$f$	$n - 1$	$n - 1$
2	2	$f - 1$	$n - 2$	$(n - 1) \cdot (n - 2)$
...	...	...	...	...
$x$	$x$	$(f + 1) - x$	$n - x$	$(n - 1)(n - 2) \dots (n - x)$
$x + 1$	$x + 1$	$(f + 1) - x - 1$	$n - x - 1$	$(n - 1)(n - 2) \dots (n - x - 1)$
$f + 1$	$f + 1$	0	$n - f - 1$	$(n - 1)(n - 2) \dots (n - f - 1)$

**Table:** Relationships between messages and rounds in the Oral Messages algorithm for Byzantine agreement.

Complexity:  $f + 1$  rounds, exponential amount of space, and

$$(n - 1) + (n - 1)(n - 2) + \dots + (n - 1)(n - 2) \dots (n - f - 1) \text{ messages}$$

## Bzantine Generals (iterative formulation), Sync, Msg-passing

(variables)

**boolean:**  $v \leftarrow$  initial value;

**integer:**  $f \leftarrow$  maximum number of malicious processes,  $\leq \lfloor \frac{n-1}{3} \rfloor$ ;

**tree of boolean:**

- level 0 root is  $v_{init}^L$ , where  $L = \langle \rangle$ ;
- level  $h$  ( $f \geq h > 0$ ) nodes: for each  $v_j^L$  at level  $h - 1 = \text{sizeof}(L)$ , its  $n - 2 - \text{sizeof}(L)$  descendants at level  $h$  are  $v_k^{\text{concat}(\langle j \rangle, L)}$ ,  $\forall k$  such that  $k \neq j$ ,  $i$  and  $k$  is not a member of list  $L$ .

(message type)

OM( $v$ ,  $Dests$ ,  $List$ ,  $faulty$ ), where the parameters are as in the recursive formulation.

(1) Initiator (i.e., Commander) initiates Oral Byzantine agreement:

(1a) **send**  $OM(v, N - \{i\}, \langle P_i \rangle, f)$  to  $N - \{i\}$ ;

```
(1b) return(v).
```

(2) (Non-initiator, i.e., Lieutenant) receives Oral Message *OM*:

```
(2a) for  $rnd = 0$  to  $f$  do
```

(2b) **for** each message OM that arrives in this round. **do**

```

(2c)   receive  $OM(v, Dests, L = \langle P_{k_1} \dots P_{k_f+1-faulty} \rangle, faulty)$  from  $P_{k_1}$ ;
        //  $faulty + round = f$ ;  $|Dests| + sizeof(L) = n$ 

```

(2d)  $v_{\text{head}(L)}^{\text{tail}(L)} \leftarrow v$ ; //  $\text{sizeof}(L) + \text{faulty} = f + 1$ . fill in estimate.

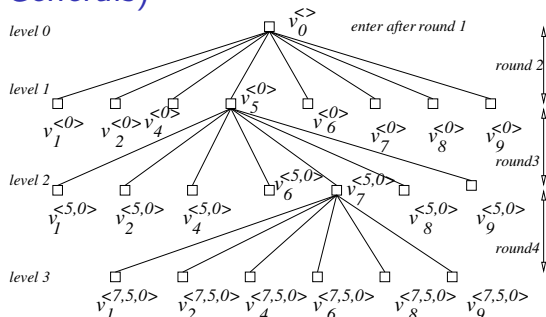
(2e) **send**  $OM(v, Dests - \{i\}, \langle P_i, P_{k_1} \dots P_{k_{f+1-faulty}} \rangle, faulty - 1)$  to  $Dests - \{i\}$  **if**  $rnd < f$ ;

```
(2f) for level = f - 1 down to 0 do
```

(2g) **for** each of the  $1 \cdot (n - 2) \cdot \dots \cdot (n - (\text{level} + 1))$  nodes  $v_Y^L$  in level *level*, **do**

$$(2h) \quad v_x^L(x \neq i, x \notin L) = \text{majority}_{y \notin \text{concat}(\langle x \rangle, L); y \neq i} (v_x^L, v_y^{\text{concat}(\langle x \rangle, L)});$$

# Tree Data Structure for Agreement Problem (Byzantine Generals)



Some branches of the tree at  $P_3$ . In

this example,  $n = 10, f = 3$ , commander is  $P_0$ .

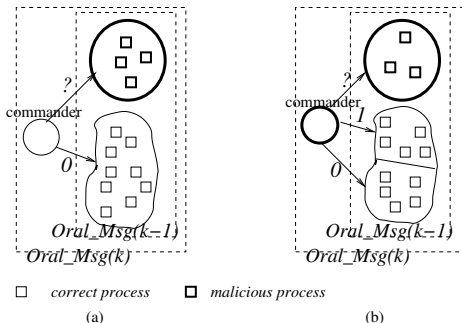
- (round 1)  $P_0$  sends its value to all other processes using  $Oral\_Msg(3)$ , including to  $P_3$ .
- (round 2)  $P_3$  sends 8 messages to others (excl.  $P_0$  and  $P_3$ ) using  $Oral\_Msg(2)$ .  $P_3$  also receives 8 messages.
- (round 3)  $P_3$  sends  $8 \times 7 = 56$  messages to all others using  $Oral\_Msg(1)$ ;  $P_3$  also receives 56 messages.
- (round 4)  $P_3$  sends  $56 \times 6 = 336$  messages to all others using  $Oral\_Msg(0)$ ;  $P_3$  also receives 336 messages. The received values are used as estimates of the majority function at this level of recursion.

# Exponential Algorithm: An example

An example of the majority computation is as follows.

- $P_3$  revises its estimate of  $v_7^{(5,0)}$  by taking  $\text{majority}(v_7^{(5,0)}, v_1^{(7,5,0)}, v_2^{(7,5,0)}, v_4^{(7,5,0)}, v_6^{(7,5,0)}, v_8^{(7,5,0)}, v_9^{(7,5,0)})$ . Similarly for the other nodes at level 2 of the tree.
- $P_3$  revises its estimate of  $v_5^{(0)}$  by taking  $\text{majority}(v_5^{(0)}, v_1^{(5,0)}, v_2^{(5,0)}, v_4^{(5,0)}, v_6^{(5,0)}, v_7^{(5,0)}, v_8^{(5,0)}, v_9^{(5,0)})$ . Similarly for the other nodes at level 1 of the tree.
- $P_3$  revises its estimate of  $v_0^{(\cdot)}$  by taking  $\text{majority}(v_0^{(\cdot)}, v_1^{(0)}, v_2^{(0)}, v_4^{(0)}, v_5^{(0)}, v_6^{(0)}, v_7^{(0)}, v_8^{(0)}, v_9^{(0)})$ . This is the consensus value.

# Impact of a Loyal and of a Disloyal Commander



effects of a loyal or a disloyal commander in a system with  $n = 14$  and  $f = 4$ . The subsystems that need to tolerate  $k$  and  $k - 1$  traitors are shown for two cases. (a) Loyal commander. (b) No assumptions about commander.


- (a) the commander who invokes  $Oral\_Msg(x)$  is loyal, so all the loyal processes have the same estimate. Although the subsystem of  $3x$  processes has  $x$  malicious processes, all the loyal processes have the same view to begin with. Even if this case repeats for each nested invocation of  $Oral\_Msg$ , even after  $x$  rounds, among the processes, the loyal processes are in a simple majority, so the majority function works in having them maintain the same common view of the loyal commander's value.
- (b) the commander who invokes  $Oral\_Msg(x)$  may be malicious and can send conflicting values to the loyal processes. The subsystem of  $3x$  processes has  $x - 1$  malicious processes, but all the loyal processes do not have the same view to begin with.



# The Phase King Algorithm

## Operation

- Each phase has a unique "phase king" derived, say, from PID.
- Each phase has two rounds:
  - 1 in 1st round, each process sends its estimate to all other processes.
  - 2 in 2nd round, the "Phase king" process arrives at an estimate based on the values it received in 1st round, and broadcasts its new estimate to all others.



# The Phase King Algorithm: Code

(variables)

**boolean:**  $v \leftarrow$  initial value;

**integer:**  $f \leftarrow$  maximum number of malicious processes,  $f < \lceil n/4 \rceil$ ;

(1) Each process executes the following  $f + 1$  phases, where  $f < n/4$ :

(1a) **for**  $phase = 1$  **to**  $f + 1$  **do**

(1b)     Execute the following Round 1 actions:                     // actions in round one of each phase

(1c)         **broadcast**  $v$  to all processes;

(1d)         **await** value  $v_j$  from each process  $P_j$ ;

(1e)          $majority \leftarrow$  the value among the  $v_j$  that occurs  $> n/2$  times (default if no maj.);

(1f)          $mult \leftarrow$  number of times that  $majority$  occurs;

(1g)     Execute the following Round 2 actions:                     // actions in round two of each phase

(1h)         **if**  $i = phase$  **then** // only the phase leader executes this send step

(1i)             **broadcast**  $majority$  to all processes;

(1j)         **receive**  $tiebreaker$  from  $P_{phase}$  (default value if nothing is received);

(1k)         **if**  $mult > n/2 + f$  **then**

(1l)              $v \leftarrow majority$ ;

(1m)         **else**  $v \leftarrow tiebreaker$ ;

(1n)         **if**  $phase = f + 1$  **then**

(1o)             output decision value  $v$ .

# The Phase King Algorithm

- $(f + 1)$  phases,  $(f + 1)[(n - 1)(n + 1)]$  messages, and can tolerate up to  $f < \lceil n/4 \rceil$  malicious processes

## Correctness Argument

- 1 Among  $f + 1$  phases, at least one phase  $k$  where phase-king is non-malicious.
- 2 In phase  $k$ , all non-malicious processes  $P_i$  and  $P_j$  will have same estimate of consensus value as  $P_k$  does.
  - 1  $P_i$  and  $P_j$  use their own majority values (Hint:  $\implies P_i$ 's *mult*  $> n/2 + f$ )
  - 2  $P_i$  uses its majority value;  $P_j$  uses phase-king's tie-breaker value. (Hint:  $P_i$ 's *mult*  $> n/2 + f$ ,  $P_j$ 's *mult*  $> n/2$  for same value)
  - 3  $P_i$  and  $P_j$  use the phase-king's tie-breaker value. (Hint: In the phase in which  $P_k$  is non-malicious, it sends same value to  $P_i$  and  $P_j$ )

In all 3 cases, argue that  $P_i$  and  $P_j$  end up with same value as estimate

- 3 If all non-malicious processes have the value  $x$  at the start of a phase, they will continue to have  $x$  as the consensus value at the end of the phase.

# Impossibility Result (MP, async)

## FLP Impossibility result

Impossible to reach consensus in an async MP system even if a single process has a crash failure

- In a failure-free async MP system, initial state is *monovalent*  $\implies$  consensus can be reached.
- In the face of failures, initial state is necessarily bivalent
- Transforming the input assignments from the all-0 case to the all-1 case, there must exist input assignments  $\vec{I}_a$  and  $\vec{I}_b$  that are 0-valent and 1-valent, resp., and that differ in the input value of only one process, say  $P_i$ . If a 1-failure tolerant consensus protocol exists, then:
  - ▶ Starting from  $\vec{I}_a$ , if  $P_i$  fails immediately, the other processes must agree on 0 due to the termination condition.
  - ▶ Starting from  $\vec{I}_b$ , if  $P_i$  fails immediately, the other processes must agree on 1 due to the termination condition.

However, execution (2) looks identical to execution (1), to all processes, and must end with a consensus value of 0, a contradiction. Hence, there must exist at least one bivalent initial state.

- Consensus requires some communication of initial values.

Key idea: in the face of a potential crash, not possible to distinguish between

# Impossibility Result (MP, async)

- To transition from bivalent to monovalent step, must exist a critical step which allows the transition by making a decision
- Critical step cannot be local (cannot tell apart between slow and failed process) nor can it be across multiple processes (it would not be well-defined)
- Hence, cannot transit from bivalent to univalent state.

## Wider Significance of Impossibility Result

- By showing reduction from consensus to problem X, then X is also not solvable under same model (single crash failure)
- E.g., leader election, terminating reliable broadcast, atomic broadcast, computing a network-wide global function using BC-CC flows, transaction commit.

# Terminating Reliable Broadcast (TRB)

A correct process always gets a message, even if sender crashes while sending (in which case the process gets a null message).

**Validity:** If the sender of a broadcast message  $m$  is non-faulty, then all correct processes eventually deliver  $m$ .

**Agreement:** If a correct process delivers a message  $m$ , then all correct processes deliver  $m$ .

**Integrity:** Each correct process delivers at most one message. Further, if it delivers a message different from the null message, then the sender must have broadcast  $m$ .

**Termination:** Every correct process eventually delivers some message.

Reduction from consensus to TRB.

- Commander sends its value using TRB.
- Receiver decides on 0 or 1 based on value it receives. If it receives a "null" message, it decides on default value.
- But, as consensus is not solvable, algo for TRB cannot exist.

# $k$ -set Consensus

**$k$ -Agreement:** All non-faulty processes must make a decision, and the set of values that the processes decide on can contain up to  $k$  ( $> f$ ) values.

**Validity:** If a non-faulty process decides on some value, then that value must have been proposed by some process.

**Termination:** Each non-faulty process must eventually decide on a value.

The  $k$ -Agreement condition is new, the Validity condition is different from that for regular consensus, and the Termination condition is unchanged from that for regular consensus.

**Example:** Let  $n = 10$ ,  $f = 2$ ,  $k = 3$  and each process choose a unique number from 1 to 10. Then 3-set is  $\{8, 9, 10\}$ .

(variables)

**integer:**  $v \leftarrow$  initial value;

(1) A process  $P_i$ ,  $1 \leq i \leq n$ , initiates  $k$ -set consensus:

(1a) **broadcast**  $v$  to all processes.

(1b) **await** values from  $|N| - f$  processes and add them to set  $V$ ;

(1c) **decide** on  $\max(V)$ .

# Epsilon Consensus (msg-passing, async)

**$\epsilon$ -Agreement:** All non-faulty processes must make a decision and the values decided upon by any two non-faulty processes must be within  $\epsilon$  range of each other.

**Validity:** If a non-faulty process  $P_i$  decides on some value  $v_i$ , then that value must be within the range of values initially proposed by the processes.

**Termination:** Each non-faulty process must eventually decide on a value.

The algorithm for the message-passing model assumes  $n \geq 5f + 1$ , although the problem is solvable for  $n > 3f + 1$ .

- Main loop simulates sync rounds.
- Main lines (1d)-(1f): processes perform all-all msg exchange
- Process broadcasts its estimate of consensus value, and awaits  $n - f$  similar msgs from other processes
- the processes' estimate of the consensus value converges at a particular rate, until it is  $\epsilon$  from any other processes estimate.
- # rounds determined by lines (1a)-(1c).



# Epsilon Consensus (msg-passing, async): Code

(variables)

**real:**  $v \leftarrow$  input value; //initial value  
**multiset of real**  $V$ ;  
**integer**  $r \leftarrow 0$ ; // number of rounds to execute

(1) Execution at process  $P_i, 1 \leq i \leq n$ :

(1a)  $V \leftarrow \text{Asynchronous\_Exchange}(v, 0)$ ;

(1b)  $v \leftarrow$  any element in( $\text{reduce}^{2f}(V)$ );

(1c)  $r \leftarrow \lceil \log_c(\text{diff}(V))/\epsilon \rceil$ , where  $c = c(n - 3f, 2f)$ .

(1d) **for round from 1 to  $r$  do**

(1e)  $V \leftarrow \text{Asynchronous\_Exchange}(v, \text{round})$ ;

(1f)  $v \leftarrow \text{new}_{2f, f}(V)$ ;

(1g) **broadcast**  $(\langle v, \text{halt} \rangle, r + 1)$ ;

(1h) **output**  $v$  as decision value.

(2)  $\text{Asynchronous\_Exchange}(v, h)$  returns  $V$ :

(2a) **broadcast**  $(v, h)$  to all processes;

(2b) **await**  $n - f$  responses belonging to round  $h$ ;


(2c) for each process  $P_k$  that sent  $\langle x, \text{halt} \rangle$  as value, use  $x$  as its input henceforth;

(2d) **return** the multiset  $V$ .

# Epsilon Consensus (msg-passing, async)

Consider a sorted collection  $U$ . The new estimate of a process is chosen by computing  $new_{k,f}(U)$ , defined as  $mean(select_k(reduce^f(U)))$

- $reduce^f(U)$  removes the  $f$  largest and  $f$  smallest members of  $U$ .
- $select_k(U)$  selects every  $k$ th member of  $U$ , beginning with the first. If  $U$  has  $m$  members,  $select_k(U)$  has  $c(m, k) = \lfloor (m - 1)/k \rfloor + 1$  members. This constant  $c$  represents a *convergence factor* towards the final agreement value, i.e., if  $x$  is the range of possible values held by correct processes before a round, then  $x/c$  is the possible range of estimate values held by those processes after that round.



*shaded members belong to  $select_5(reduce^4(U))$*

$select_k(reduce^f(U))$  operation, with  $k = 5$  and  $f = 4$ . The mean of the selected members is the new estimate  $new_{5,4}(U)$ .

The algorithm uses  $m = n - 3f$  and  $k = 2f$ . So  $c(n - 3f, 2f)$  will represent the *convergence factor* towards reaching approximate agreement and  $new_{2f,f}$  is the new estimate after each round.

# Epsilon Consensus (msg-passing, async)

Let  $|U| = m$ , and let the  $m$  elements  $u_0 \dots u_{m-1}$  of multiset  $U$  be in nondecreasing order.


Properties on nonempty multisets  $U, V, W$ .

- The number of the elements in multisets  $U$  and  $V$  is reduced by at most 1 when the smallest element is removed from both. Similarly for the largest element.
- The number of elements common to  $U$  and  $V$  before and after  $j$  reductions differ by at most  $2j$ . Thus, for  $j \geq 0$  and  $|V|, |W| \geq 2j$ ,  
 $|V \cap W| - |\text{reduce}^j(V) \cap \text{reduce}^j(W)| \leq 2j$ .
- Let  $V$  contain at most  $j$  values not in  $U$ , i.e.,  $|V - U| \leq j$ , and let size of  $V$  be at least  $2j$ . Then by removing the  $j$  low and  $j$  high elements from  $V$ , it is easy to see that remaining elements in  $V$  must belong to the range of  $U$ .

Thus,

- each value in  $\text{reduce}^j(V)$  is in the range of  $U$ , i.e.,  $\text{range}(\text{reduce}^j(V)) \subseteq \text{range}(U)$ .
- $\text{new}_{k,j}(V) \in \text{range}(U)$ .

Correctness, termination, complexity: refer book



# Asynchronous Renaming

The renaming problem assigns to each process  $P_i$ , a name  $m_i$  from a domain  $M$ , and is formally specified as follows.

- Agreement:** For non-faulty processes  $P_i$  and  $P_j$ ,  $m_i \neq m_j$ .
- Termination:** Each non-faulty process is eventually assigned a name  $m_i$ .
- Validity:** The name  $m_i$  belongs to  $M$ .
- Anonymity:** The code executed by any process must not depend on its initial identifier.

Uses of renaming (name space transformation):

- processes from different domains need to collaborate, but must first assign themselves distinct names from a small domain.
- processes need to use their names as “tags” to simply mark their presence, as in a priority queue.
- the name space has to be condensed, e.g., for  $k$ -mutex.

Assumptions

- The  $n$  processes  $P_1 \dots P_n$  have their identifiers in the old name space.  $P_i$  knows only its identifier, and the total number of processes,  $n$ .
- The  $n$  processes take on new identifiers  $m_1 \dots m_n$ , resp., from the name space  $M$ .
- Due to asynchrony, each process that chooses its new name must continue to cooperate with the others until they have chosen their new names.

# Asynchronous Renaming -MP Model

- Attiya et al. renaming algorithm assumes  $n \geq 2f + 1$  and fail-stop model.
- Transformed name space is  $M = n + f$ .
- *View* is a list of up to  $n$  objects of type *bid*.

(local variables)

**struct** *bid*:

```

integer P;                                // old name of process
integer x;                                // new name being bid by the process
integer attempt;                          // the number of bids so far, including this current bid
boolean decide;                           // whether new name x is finalized

```

```

list of bid: View[1 . . . n]  $\leftarrow \langle \langle i, 0, 0, false \rangle \rangle$ ; // initialize list with an entry for  $P_i$ 

```

```

integer count;                            // number of copies of the latest local view, received from others

```

```

boolean: restart, stable, no_choose;      // loop control variables

```

- $View \leq View'$  if and only if for each process  $P_i$  such that  $View[k].P = P_i$ , we also have that for some  $k'$ ,  $View'[k'].P = P_i$  and  $View[k].attempt \leq View'[k'].attempt$ .

If  $View' \not\leq View$  (line 1n), then *View* is updated using *View'* (line 1o) by:

- 1 including all process entries from *View'* that are missing in *View* (i.e.,  $View'[k'].P$  is not equal to  $View[k].P$ , for all  $k$ ), so such entries  $View'[k']$  are added to *View*.
- 2 replacing older entries for the same process with more recent ones, (i.e., if  $View'[k'].P = P_i = View[k].P$  and  $View'[k'].attempt > View[k].attempt$ , replace  $View[k]$  by  $View'[k']$ ).



# Wait-free Renaming: Code

(1) A process  $P_i$ ,  $1 \leq i \leq n$ , participates in renaming:

```

(1a) repeat
(1b)    $restart \leftarrow false$ ;
(1c)   broadcast  $message(View)$ ;
(1d)    $count \leftarrow 1$ ;
(1e)   repeat
(1f)      $no\_choose \leftarrow 0$ ;
(1g)     repeat
(1h)       await  $message(View')$ ;
(1i)        $stable \leftarrow false$ ;
(1j)       if  $View' = View$  then
(1k)          $count \leftarrow count + 1$ ;
(1l)         if  $count \geq n - f$  then
(1m)            $stable \leftarrow true$ ;
(1n)       else if  $View' \not\leq View$  then
(1o)         update  $View$  using  $View'$  by taking latest information for each process;
(1p)          $restart \leftarrow true$ ;
(1q)   until ( $stable = true$  or  $restart = true$ ); //  $n - f$  copies received, or new view obtained
(1r)   if  $restart = false$  then //  $View[1]$  has information about  $P_i$ 
(1s)     if  $View[1].x \neq 0$  and  $View[1].x \neq View[j].x$  for any  $j$  then
(1t)       decide  $View[1].x$ ;
(1u)        $View[1].decide \leftarrow true$ ;
(1v)       broadcast  $message(View)$ ;
(1w)   else
(1x)     let  $r$  be the rank of  $P_i$  in  $UNDECIDED(View)$ ;
(1y)     if  $r \leq f + 1$  then
(1z)        $View[1].x \leftarrow FREE(View)(r)$ , the  $r$ th free name in  $View$ ;
(1A)        $View[1].attempt \leftarrow View[1].attempt + 1$ ;
(1B)        $restart \leftarrow 1$ ;
(1C)   else
(1D)      $no\_choose \leftarrow 1$ ;
(1E)   until  $no\_choose = 0$ ;
(1F)   until  $restart = 0$ ;
(1G)   repeat
(1H)     on receiving  $message(View')$ 
(1I)       update  $View$  with  $View'$  if necessary;
(1J)       broadcast  $message(View)$ ;
(1K)   until  $false$ .

```

# Reliable Broadcast


- Reliable Broadcast is RTB without terminating condition.
- RTB requires eventual delivery of messages, even if sender fails before sending. In this case, a null message needs to get sent. In RB, this condition is not there.
- RTB requires recognition of a failure, even if no msg is sent
- Crux: RTB is required to distinguish between a failed process and a slow process.
- RB is solvable under crash failures;  $O(n^2)$  messages

- (1) Process  $P_0$  initiates Reliable Broadcast:
  - (1a) **broadcast** message  $M$  to all processes.
- (2) A process  $P_i$ ,  $1 \leq i \leq n$ , receives message  $M$ :
  - (2a) **if**  $M$  was not received earlier **then**
  - (2b)     **broadcast**  $M$  to all processes;
  - (2c)     deliver  $M$  to the application.




# Shared Memory Consensus (async): Impossibility

- Use FLP argument seen in async MP systems here for SM systems.
- Cannot distinguish between failed process and a slow process  $\implies$  consensus not possible.
- Proof by contradiction, using notion of *critical step* at which system transitions from bivalent to monovalent state.
- Given initial bivalent state, prefix  $Z$ , then step by  $P_i$  leads to 0-valent state but event at some  $P_j$  followed by step of  $P_i$  leads to 1-valent state.
- Apply case analysis on prefix  $Z$  and actions of  $P_i$  and  $P_j$  after  $Z$ .




# Shared Memory Consensus (async): Impossibility


- (a)  $P_i$  does a Read.  $extend(Z, i \cdot j)$  and  $extend(Z, j \cdot i)$  are isomorphic to all except  $P_i$ . If  $P_i$  stops after  $extend(Z, i \cdot j)$ , all must reach consensus 0 after some suffix  $\delta$ . However, as per Figure (a), processes must reach consensus 1 after  $\delta$ . A contradiction.
- (a')  $P_j$  does a Read. Similar reasoning to case (a)
- (b)  $P_i$  and  $P_j$  Write to different vars. System state after  $extend(Z, i \cdot j)$  and  $extend(Z, j \cdot i)$  will have to be 0-valent and 1-valent, resp.. A contradiction.
- (c)  $P_i$  and  $P_j$  Write to the same variable. System states after  $extend(Z, i)$  and after  $extend(Z, j \cdot i)$  are isomorphic to all except  $P_j$ . Assume  $P_j$  does not run now. Then a contradiction can be seen, because of consensus value 0 after the first prefix and a consensus value of 1 after the second prefix.



(a)  $i$  does a Read  
(same logic if  
 $j$  does a Read)



(b)  $i$  and  $j$  write to  
different variables



(c)  $i$  and  $j$  write to  
the same variable

# Wait-free SM Consensus using Shared Objects

Not possible to go from bivalent to univalent state if even a single failure is allowed. Difficulty is not being able to read & write a variable atomically.

- It is not possible to reach consensus in an asynchronous shared memory system using Read/Write atomic registers, even if a single process can fail by crashing.
- There is no wait-free consensus algorithm for reaching consensus in an asynchronous shared memory system using Read/Write atomic registers.

To overcome these negative results

- Weakening the consensus problem, e.g.,  $k$ -set consensus, approximate consensus, and renaming using atomic registers.
- Using memory that is stronger than atomic Read/Write memory to design wait-free consensus algorithms. Such a memory would need corresponding access primitives.

## Stronger objects?

Are there objects (with supporting operations), using which there is a wait-free (i.e.,  $(n - 1)$ -crash resilient) algorithm for reaching consensus in a  $n$ -process system? Yes, e.g., Test&Set, Swap, Compare&Swap.

Henceforth, assume only the crash failure model, and also require the solutions to be *wait-free*.

# Consensus Numbers and Consensus Hierarchy

## Consensus Numbers

An object of type  $X$  has consensus number  $k$ , denoted as  $CN(X) = k$ , if  $k$  is the largest number for which the object  $X$  can solve wait-free  $k$ -process consensus in an asynchronous system subject to  $k - 1$  crash failures, using only objects of type  $X$  and read/write objects.

## Wait-free simulations and Consensus Numbers

For objects  $X$  and  $Y$  such that  $CN(X) < CN(Y)$ , there is no wait-free simulation of object  $Y$  using  $X$  and read/write registers (whose consensus number is 1) in a system with more than  $CN(X)$  processes.

There does not exist any simulation of objects with  $CN > 1$  using only Read/Write atomic registers  $\implies$  need stronger objects.

Object	Consensus number
Read/Write objects	1
Test-&-Set, stack, FIFO queue, Fetch-&-Inc	2
Augmented queue with peek - size $k$	$k$
Compare-&-Swap, Augmented queue, memory-memory move memory-memory swap, Fetch-&-Cons, store-conditional	$\infty$

# Definitions of Sync Operations *RMW*, Compare&Swap, Fetch&Inc

(shared variables among the processes accessing each of the different object types)

**register:**  $Reg \leftarrow$  initial value;

// shared register initialized

(local variables)

**integer:**  $old \leftarrow$  initial value;

// value to be returned

**integer:**  $key \leftarrow$  comparison value for conditional update;

(1)  $RMW(Reg, \text{function } f)$  returns  $value$ :

(1a)  $old \leftarrow Reg$ ;

(1b)  $Reg \leftarrow f(Reg)$ ;

(1c) **return**( $old$ ).

(2)  $Compare\&Swap(Reg, key, new)$  returns  $value$ :

(2a)  $old \leftarrow Reg$ ;

(2b) **if**  $key = old$  **then**

(2c)      $Reg \leftarrow new$ ;

(2d) **return**( $old$ ).

(3)  $Fetch\&Inc(Reg)$  returns  $value$ :

(3a)  $old \leftarrow Reg$ ;

(3b)  $Reg \leftarrow r + 1$ ;

(3c) **return**( $old$ ).

# Two-process Wait-free Consensus using FIFO Queue

```

(shared variables)
queue:  $Q \leftarrow \langle 0 \rangle;$                                      // queue  $Q$  initialized
integer:  $Choice[0, 1] \leftarrow [\perp, \perp]$              // preferred value of each process
(local variables)
integer:  $temp \leftarrow 0;$ 
integer:  $x \leftarrow \text{initial choice};$ 

```

(1) Process  $P_i, 0 \leq i \leq 1$ , executes this for 2-process consensus using a FIFO queue:

```

(1a)  $Choice[i] \leftarrow x;$ 
(1b)  $temp \leftarrow \text{dequeue}(Q);$ 
(1c) if  $temp = 0$  then
(1d)   output( $x$ )
(1e) else output( $Choice[1 - i]$ ).

```

# Wait-free Consensus using Compare&Swap

(shared variables)

**integer:**  $Reg \leftarrow \perp$ ;

// shared register  $Reg$  initialized

(local variables)

**integer:**  $temp \leftarrow 0$ ;

// temp variable to read value of  $Reg$

**integer:**  $x \leftarrow$  initial choice;

// initial preference of process

(1) Process  $P_i$ , ( $\forall i \geq 1$ ), executes this for consensus using *Compare&Swap*:

(1a)  $temp \leftarrow \text{Compare\&Swap}(Reg, \perp, x)$ ;

(1b) **if**  $temp = \perp$  **then**

(1c)     **output**( $x$ )

(1d) **else output**( $temp$ ).

# Read-Modify-Write (MRW) Abstraction

- RMW allows to read, and modify the register content as per some function  $f$ .
- RMW object has a CN of at least 2 because it allows the first process to access the object to leave an imprint that the object has been accessed. The other process can read the imprint.
- If the imprint can include the ID of the first process, or the choice of the first process, then  $CN > 2$ .
- RMW objects differ in their function  $f$ . A function is termed as *interfering* if for all process pairs  $i$  and  $j$ , and for all legal values  $v$  of the register,
  - 1  $f_i(f_j(v)) = f_j(f_i(v))$ , i.e., function is commutative, or
  - 2 the function is not write-preserving, i.e.,  $f_i(f_j(v)) = f_i(v)$  or vice-versa with the roles of  $i$  and  $j$  interchanged.
- **Examples:**
  - ▶ The *Fetch&Inc* commutes even though it is write-preserving.
  - ▶ The *Test&Set* commutes and is not write-preserving.
  - ▶ The *Swap* does not commute but it is not write-preserving.

Hence, all three objects uses functions that are *interfering*.



# RMW Object and Instruction

A nontrivial interfering RMW operation has consensus number = 2

- If RMW is commutative, 3rd process cannot know which of the other two accessed the object first, and therefore does not know whose value is the consensus value
- If RMW is not write-preserving, 3rd process does not know if it is the 2nd or 3rd to access the object. Therefore, whose value is the consensus value?

Objects like Compare&Swap are non-interfering and hence have a higher consensus number.

# RMW Object and Instruction

(shared variables)

**integer:**  $Reg \leftarrow \perp$ ;

// shared register  $Reg$  initialized

**integer:**  $Choice[0, 1] \leftarrow [\perp, \perp]$ ;

// data structure

(local variables)

**integer:**  $x \leftarrow \text{initial choice}$ ;

// initial preference of process

(1) Process  $P_i$ , ( $0 \leq i \leq 1$ ), executes this for consensus using *RMW*:

(1a)  $Choice[i] \leftarrow x$ ;

(1b)  $val \leftarrow RMW(Reg, f)$ ;

(1c) **if**  $val = \perp$  **then**

(1d)           **output**( $Choice[i]$ )

(1e) **else output**( $Choice[1 - i]$ ).

**RMW register**

$Reg$



$Choice$   $[0]$   $[1]$

# Universality of Consensus Objects


- An object is defined to be *universal* if that object along with read/write registers can simulate any other object in a wait-free manner. In any system containing up to  $k$  processes, an object  $X$  such that  $CN(X) = k$  is *universal*.
- For any system with up to  $k$  processes, the universality of objects  $X$  with consensus number  $k$  is shown by giving a *universal* algorithm to wait-free simulate *any* object using only objects of type  $X$  and read/write registers. This is shown in two steps.
  - 1 A *universal* algorithm to wait-free simulate *any* object whatsoever using read/write registers and arbitrary  $k$ -processor consensus objects is given. This is the main step.
  - 2 Then, the arbitrary  $k$ -process consensus objects are simulated with objects of type  $X$ , also having consensus number  $k$ . This trivially follows after the first step.
- Hence, any object  $X$  with consensus number  $k$  is universal in a system with  $n \leq k$  processes.

# Universality of Consensus Objects

- An arbitrary consensus object  $X$  allows a single operation,  $Decide(X, v_{in})$  and returns a value  $v_{out}$ , where both  $v_{in}$  and  $v_{out}$  have to assume a legal value from known domains  $V_{in}$  and  $V_{out}$ , resp.
- For the correctness of this shared object version of the consensus problem, all  $v_{out}$  values returned to each invoking process must equal the  $v_{in}$  of some process.
- A *nonblocking* operation, in the context of shared memory operations, is an operation that may not complete itself but is guaranteed to complete at least one of the pending operations in a finite number of steps.

# A Nonblocking Universal Algorithm

- The linked list stores the linearized sequence of operations and states following each operation.
- Operations to the arbitrary object  $Z$  are simulated in a nonblocking way using only an arbitrary consensus object (namely, the field  $op.next$  in each record) which is accessed via the *Decide* call.
- Each process attempts to thread its own operation next into the linked list.



# A Nonblocking Universal Algorithm

(shared variables)

**record** *op*

**integer:** *seq*  $\leftarrow 0$ ;

*operation*  $\leftarrow \perp$ ;

*state*  $\leftarrow$  *initial state*;

*result*  $\leftarrow \perp$ ;

**op** *\*next*  $\leftarrow \perp$ ;

// sequence number of serialized operation

// operation, with associated parameters

// the state of the object after the operation

// the result of the operation, to be returned to invoker

// pointer to the next record

**op**  $*Head[1 \dots k] \leftarrow \&(anchor\_record)$ ;

(local variables)

**op**  $*my\_new\_record, *winner$ ;

(1) Process  $P_i, 1 \leq i \leq k$  performs operation *invoc* on an arbitrary consensus object:

(1a)  $my\_new\_record \leftarrow malloc(op)$ ;

(1b)  $my\_new\_rec.operation \leftarrow invoc$ ;

(1c) **for** *count* = 1 **to** *k* **do**

(1d)     **if**  $Head[i].seq < Head[count].seq$  **then**

(1e)          $Head[i] \leftarrow Head[count]$ ;

(1f) **repeat**

(1g)      $winner \leftarrow Decide(Head[i].next, \&my\_new\_record)$ ;

(1h)      $winner.seq \leftarrow Head[i].seq + 1$ ;

(1i)      $winner.state, winner.result \leftarrow apply(winner.operation, Head[i].state)$ ;


(1j)      $Head[i] \leftarrow winner$ ;

(1k) **until**  $winner = my\_new\_record$ ;

(1l) enable the response to *invoc*, that is stored at *winner.result*.

# A Nonblocking Universal Algorithm: Notes

- There are as many universal objects as there are operations to thread.
- A single pointer/counter cannot be used instead of the array *Head*. B'coz reading and updating the pointer cannot be done atomically in a wait-free manner.
- Linearization of the operations given by the seq no.
- As algorithm is nonblocking, some process(es) may be starved indefinitely.



# A Wait-free Universal Algorithm

(shared variables)

**record** *op*

**integer:** *seq*  $\leftarrow$  0;

*operation*  $\leftarrow \perp$ ;

*state*  $\leftarrow$  *initial state*;

*result*  $\leftarrow \perp$ ;

**op** *\*next*  $\leftarrow \perp$ ;

// sequence number of serialized operation

// operation, with associated parameters

// the state of the object after the operation

// the result of the operation, to be returned to invoker

// pointer to the next record

**op** *\*Head*[1 .. *k*], *\*Announce*[1 .. *k*]  $\leftarrow \overline{\&(\text{anchor\_record})}$ ;

(local variables)

**op** *\*my\_new\_record*, *\*winner*;

(1) Process  $P_i$ ,  $1 \leq i \leq k$  performs operation *invoc* on an arbitrary consensus object:

(1a) *Announce*[*i*]  $\leftarrow$  *malloc*(*op*);

(1b) *Announce*[*i*].*operation*  $\leftarrow$  *invoc*; *Announce*[*i*].*seq*  $\leftarrow$  0;

(1c) **for** *count* = 1 **to** *k* **do**

(1d) **if** *Head*[*i*].*seq* < *Head*[*count*].*seq* **then**

(1e) *Head*[*i*]  $\leftarrow$  *Head*[*count*];

(1f) **while** *Announce*[*i*].*seq* = 0 **do**

(1g) *turn*  $\leftarrow$  (*Head*[*i*].*seq* + 1) mod (*k*);

(1h) **if** *Announce*[*turn*].*seq* = 0 **then**

(1i) *my\_new\_record*  $\leftarrow$  *Announce*[*turn*];

(1j) **else** *my\_new\_record*  $\leftarrow$  *Announce*[*i*];

(1k) *winner*  $\leftarrow$  *Decide*(*Head*[*i*].*next*, &*my\_new\_record*);

(1l) *winner*.*seq*  $\leftarrow$  *Head*[*i*].*seq* + 1;

(1m) *winner*.*state*, *winner*.*result*  $\leftarrow$  *apply*(*winner*.*operation*, *Head*[*i*].*state*);


(1n) *Head*[*i*]  $\leftarrow$  *winner*;

(1o) enable the response to *invoc*, that is stored at *winner*.*result*.



# Wait-free Universal Algorithm

- To prevent starvation in the nonblocking algorithm, the idea of "helping" using a round-robin approach modulo  $n$  is used.
- If  $P_j$  determines that the next op is to be assigned sequence number  $x$ , then it first checks whether the process  $P_i$  such that  $i = x \pmod{n}$  is contending for threading its operation. If so, then  $P_j$  tries to thread  $P_i$ 's operation instead of its own.
- The round-robin approach uses the array *Announce*.
- Within  $n$  iterations of the outer loop, a process is certain that its operation gets threaded - by itself or with the help of another contending process.



# Shared Memory $k$ -set Consensus

- Crash failure model,  $k > f$ . Analogous to message-passing model algorithm. Assumes atomic snapshot object  $Obj$ .
- $P_i$  writes its value to  $Obj[i]$  and scans  $Obj$  until  $n - f$  values have been written to it. Then takes the max.

(variables)

**integer:**  $v \leftarrow$  initial value;

**array of integer**  $local\_array \leftarrow \perp$ ;

(shared variables)

**atomic snapshot object**  $Obj[1 \dots n] \leftarrow \perp$ ;

(1) A process  $P_i, 1 \leq i \leq n$ , initiates  $k$ -set consensus:

(1a) **update** $_i(Obj[i])$  with  $v$ ;

(1b) **repeat**

(1c)  $local\_array \leftarrow \text{scan}_i(Obj)$ ;

(1d) **until** there are at least  $|N| - f$  non-null values in  $Obj$ ;

(1e)  $v \leftarrow$  max. of the values in  $local\_array$ .

# Async Wait-free Renaming using Atomic Shared Object

- Crash failure model. *Obj* linearizes all accesses to it.
- Each  $P_i$  can write to its portion in *Obj* and read all *Obj* atomically.
- $P_i$  does not have a unique index from  $[1 \dots n]$ .
- $P_i$  proposes a name "1" for itself. It then repeats the following loop.
  - ▶ It writes its latest bid to its component of *Obj* (line 1c); it reads the entire object using a **scan** into its local array (line 1d).  $P_i$  examines the local array for a possible conflict with its proposed new name (line 1e).
    - ★ If  $P_i$  detects a conflict with its proposed name  $m_i$  (line 1e) it determines its rank  $rank$  among the *old* names (line 1f); and selects the  $rank^{th}$  smallest integer among the names that have not been proposed in the view just read (line 1g). This will be used as  $P_i$ 's bid for a new name in the next iteration.
    - ★ If  $P_i$  detects no conflict with its proposed name  $m_i$  (line 1e), it selects this name and exits (line 1i).

# Async Wait-free Renaming using Atomic Shared Object

**Correctness:** As  $Obj$  is linearizable, no two processes having chosen a new name will get back a Scan saying their new names are unique.

**Size of new name space:**  $[1 \dots 2n - 1]$ .

**Termination:** Assume there is a subset  $\overline{T} \subseteq N$  of processes that never terminate. Let  $\min(\overline{T})$  be the process in  $\overline{T}$  with the lowest ranked process identifier (old name). Let  $\text{rank}(\min(\overline{T}))$  be the rank of this process among *all* the processes  $P_1 \dots P_n$ . Once every process in  $\overline{T}$  has done at least one **update**, and once all the processes in  $T$  have terminated, we have the following.

- The set of names of the terminated processes, say  $M_T$ , remains fixed.
- The process  $\min(\overline{T})$  will choose a name not in  $M_T$ , that is ranked  $\text{rank}(\min(\overline{T}))$ . As  $\text{rank}(\min(\overline{T}))$  is unique, no other process in  $\overline{T}$  will ever choose this name.
- Hence,  $\min(\overline{T})$  will not detect any conflict with  $\text{rank}(\min(\overline{T}))$  and will terminate.

As  $\min(\overline{T})$  cannot exist, the set  $\overline{T} = \emptyset$ .

**Lower bound:** For crash-failures, lower bound of  $n + f$  on new name space.

# Async Wait-free Renaming using Atomic Shared Object

(variables)

**integer:**  $m_i \leftarrow 0$ ;

**integer:**  $P_i \leftarrow$  name from old domain space;

**list of integer tuples**  $local\_array \leftarrow \langle \perp, \perp \rangle$ ;

(shared variables)

**atomic snapshot object**  $Obj \leftarrow \overline{\langle \perp, \perp \rangle}$ ; //  $n$  components

(1) A process  $P_i, 1 \leq i \leq n$ , participates in wait-free renaming:

(1a)  $m_i \leftarrow 1$ ;

(1b) **repeat**

(1c)     **update** $_i(Obj, \langle P_i, m_i \rangle)$ ; // update  $i$ th component with bid  $m_i$

(1d)      $local\_array(\langle P_1, m_1 \rangle, \dots, \langle P_n, m_n \rangle) \leftarrow \text{scan}_i(Obj)$ ;

(1e)     **if**  $m_i = m_j$  for some  $j \neq i$  **then**

(1f)         Determine rank  $rank_i$  of  $P_i$  in  $\{P_j \mid P_j \neq \perp \wedge j \in [1, n]\}$ ;

(1g)          $m_k \leftarrow rank_i$ th smallest integer not in  $\{m_j \mid m_j \neq \perp \wedge j \in [1, n] \wedge j \neq i\}$ ;

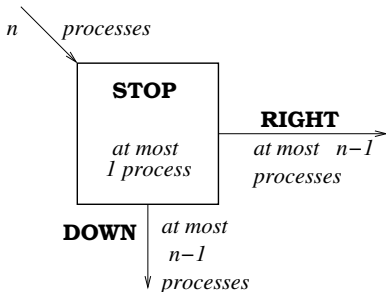
(1h)     **else**

(1i)         **decide** $(m_k)$ ; **exit**;

(1j) **until** *false*.

# The Splitter

- At most one process is returned *stop*.
- At most  $n - 1$  processes are returned *down*.
- At most  $n - 1$  processes are returned *right*.



(shared variables)

**MRMW atomic snapshot object**  $X$ ,  $\leftarrow 0$ ; **MRMW atomic snapshot object**  $Y \leftarrow false$ ;

(1) *splitter()*, executed by process  $P_i, 1 \leq i \leq n$ :

(1a)  $X \leftarrow i$ ;

(1b) **if**  $Y$  **then**

(1c)     **return**(*right*);


(1d) **else**

(1e)      $Y \leftarrow true$ ;

(1f)     **if**  $X = i$  **then** **return**(*stop*)

(1g)     **else** **return**(*down*).

# Configuration of Splitters for Wait-free Renaming (SM)



name space:  $n(n+1)/2$  splitters

(local variables)

$next, r, d, new\_name \leftarrow 0;$

(1) Process  $P_i, 1 \leq i \leq n$ , participates in wait-free renaming:

(1a)  $r, d \leftarrow 0;$

(1b) **while**  $next_i \neq stop$  **do**

(1c)  $next_i \leftarrow splitter(r, d);$

(1d) **case**

(1e)  $next = right$  **then**  $r \leftarrow r + 1;$

(1f)  $next = down$  **then**  $d \leftarrow d + 1;$

(1g)  $next = stop$  **then break()**

(1h) **return**( $new\_name = n \cdot d - d(d-1)/2 + r$ ).

New