

## Unit II -Programming Shared Address Space Platforms:

---

- ✓ Thread
  - ✓ Basics
  - ✓ The POSIX Thread API,
  - ✓ Thread Creation and Termination,
  - ✓ Synchronization Primitives in Pthreads,
  - ✓ Controlling Thread and Synchronization Attributes,
  - ✓ Thread Cancellation, Composite Synchronization Constructs.

# Distributed memory systems

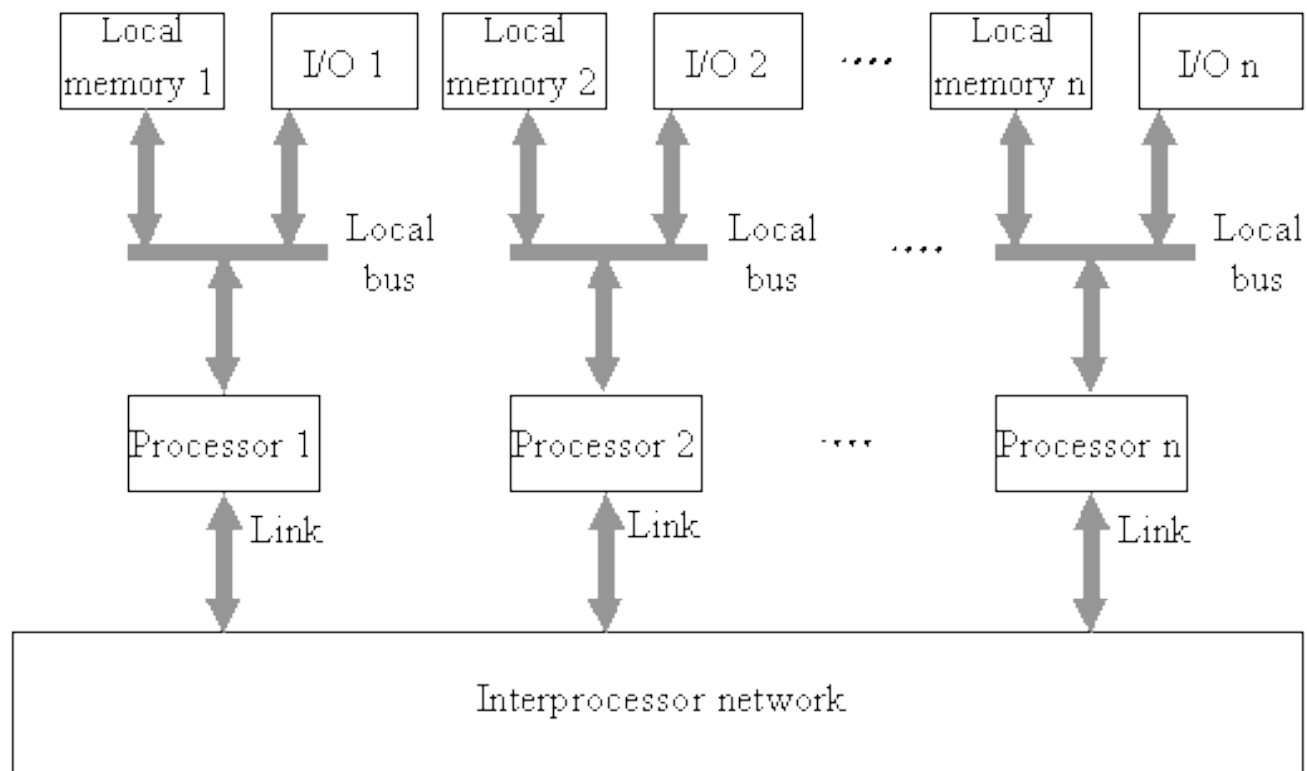
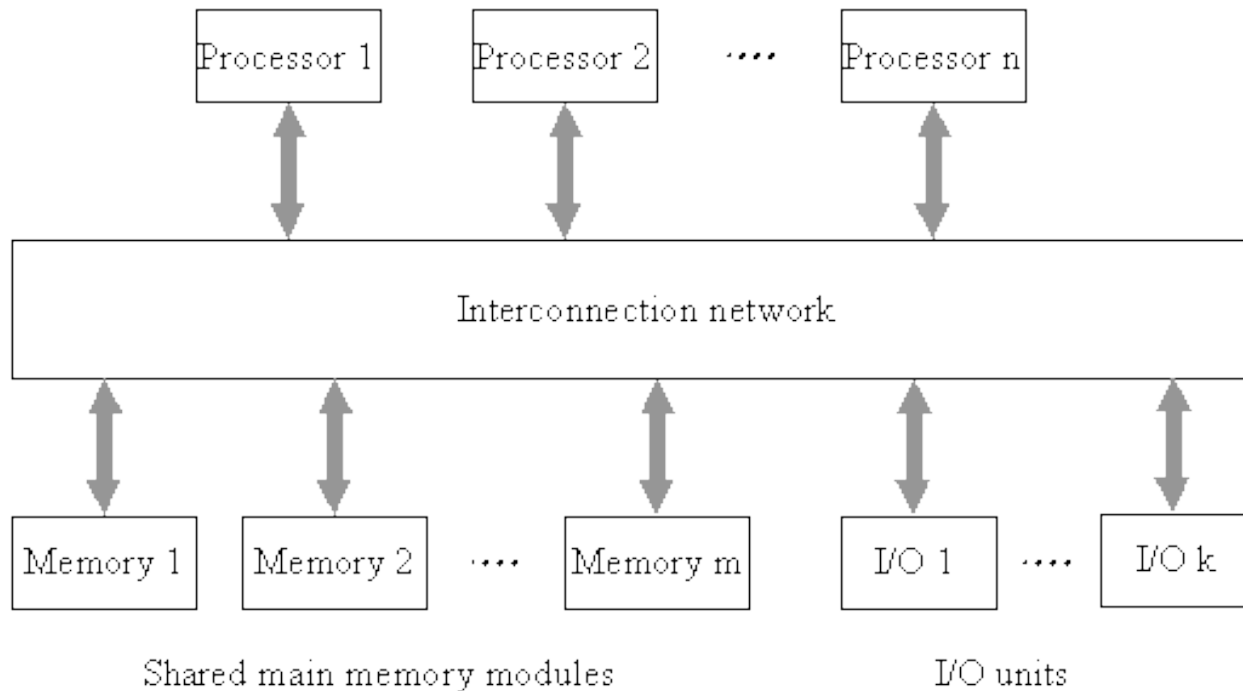


Fig: A multiprocessor system with a distributed memory (loosely coupled system)

## Distributed memory

- Each processor has its own private memory.
- Computational tasks can only operate on local data,
- if remote data is required, the computational task must communicate with one or more remote processors. Communication through the **message passing**.

# Shared memory systems




all processors can access all the main memory address space.

Shared variables access in the main memory


Fetching instructions for execution in processors is also done from a shared memory

A multiprocessor system with shared memory (tightly coupled system)

## Multiprocessing



## Multithreading



# Threaded Programming Models

- Library-based models —
  - all data is shared, unless otherwise specified
  - Examples: Pthreads, Intel Threading Building Blocks, Java Concurrency, Boost, Microsoft .Net Task Parallel Library •
- Directive-based models
  - e.g., OpenMP —shared and private data —
  - Thread creation and synchronization
  - Programming languages —
    - Cilk Plus (Intel, GCC) —CUDA (NVIDIA) —Habanero-Java (Rice/Georgia Tech)

# Parallel Programming

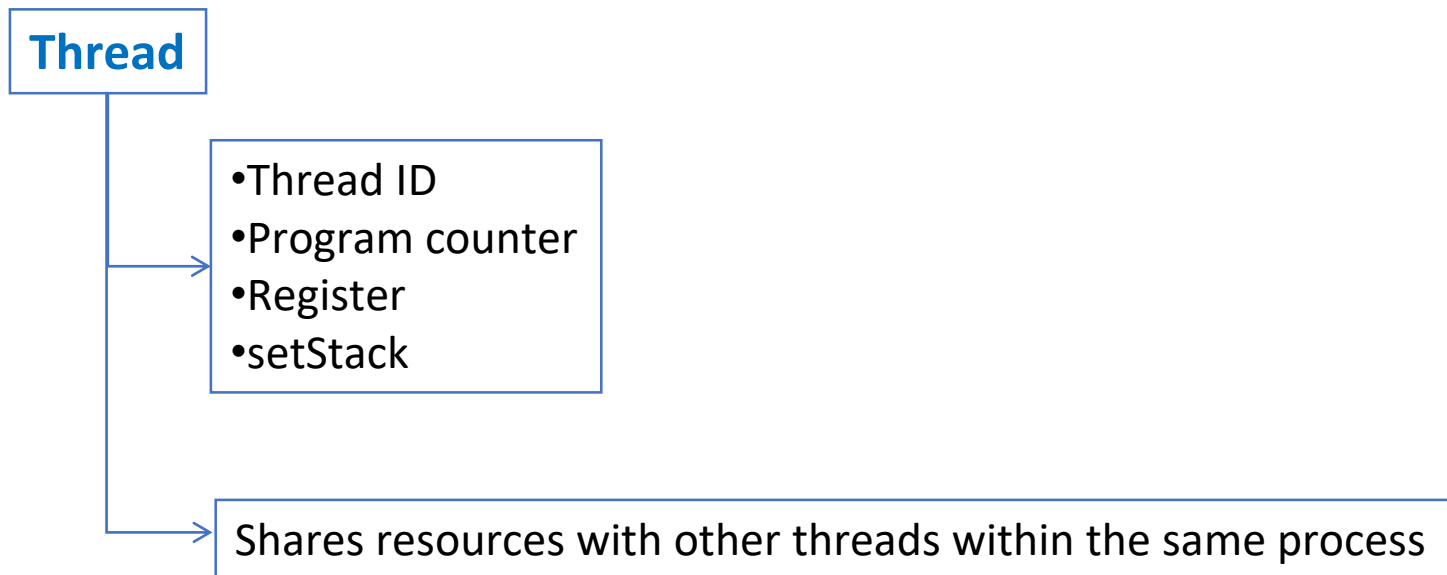
1. Synchronization between concurrent tasks
2. Communication of intermediate results

## Shared address space architectures -



- Implicitly specified since some (or all) of the memory is accessible to all the processors.
- **Threads** assume that all memory is global

# Thread Basics

- ***A thread is a single stream of control in the flow of a program and it is a basic unit of CPU utilization.***



# Thread Basics



- The logical machine model of a thread-based programming paradigm.



# Why Threads?

- Threads provide **software portability**.
- Inherent support **for latency hiding**.
- **Scheduling and load balancing**.
- Ease **of programming and widespread use**.

## Product of two dense matrices of size $n \times n$ .

- ```
for (row = 0; row < n; row++)  
    for (column = 0; column < n; column++)  
        c[row][column] =  
            dot_product( get_row(a, row),  
                        get_col(b, col));
```

**can be transformed to:**

```
for (row = 0; row < n; row++)  
    for (column = 0; column < n; column++)  
        c[row][column] =  
            create_thread( dot_product(get_row(a,  
row),  
col))) );
```

# The POSIX Thread API

- Number of vendors provide vendor-specific thread APIs
- Pthreads-IEEE specifies a standard 1003.1c-1995, **POSIX API**
- POSIX has emerged as the standard threads API
- Other thread API: NT threads, Solaris threads, Java threads, etc

# Thread Basics: Creation and Termination

- Pthreads provides two basic functions for specifying concurrency in a program:

```
#include <pthread.h>

int pthread_create (
    pthread_t *thread_handle, const pthread_attr_t *attribute,
    void * (*thread_function) (void *),
    void *arg);

int pthread_join (
    pthread_t thread,
    void **ptr);
```

- The function `pthread_create` invokes function `thread_function` as a thread

# Create thread


- `int pthread_create( pthread_t *thread, pthread_attr_t *attr, void *(*thread_function)(void *), void *arg );`
- 1st arg – pointer to the identifier of the created thread
- 2nd arg – thread attributes. If null, then the thread is created with default attributes
- 3rd arg – pointer to the function the thread will execute
- 4th arg – the argument of the executed function
- returns 0 for success

# Waiting threads

`int pthread_join( pthread_t thread, void **thread_return )`

- main thread will wait for daughter thread *thread* to finish
- 1st arg – the thread to wait for
- 2nd arg – pointer to a pointer to the return value from the thread
- returns 0 for success
- threads should always be joined; otherwise, a thread might keep on running even when the main thread has already terminated

# Estimating Pi using the Monte Carlo Method




The area of the circle is  $\pi r^2$ ,  
The area of the square is  $\text{width}^2 = (2r)^2 = 4r^2$ .  
If we divide the area of the circle, by the area of the square we get  $\pi/4$ .

## Running Example: Monte Carlo Estimation of Pi

### Approximate Pi

- generate random points with  $x, y \in [-0.5, 0.5]$
- test if point inside the circle, i.e.,  
 $x^2 + y^2 < (0.5)^2$
- ratio of circle to square =  
 $\pi r^2 / 4r^2 = \pi / 4$
- $\pi \approx 4 * (\text{number of points inside the circle}) / (\text{number of points total})$



# Thread Basics: Creation and Termination (Example)

```
#include <pthread.h>
#include <stdlib.h>
#define MAX_THREADS 512
void *compute_pi (void *);
....
main() {
    ...
    pthread_t p_threads[MAX_THREADS];
    pthread_attr_t attr;
    pthread_attr_init (&attr);
    for (i=0; i< num_threads; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void *) &hits[i]);
    }
    for (i=0; i< num_threads; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
    ...
}
```



# Thread Basics: Creation and Termination (Example)

```
#include <pthread.h>
#include <stdlib.h>
#define MAX_THREADS 512
void *compute_pi (void *);
....
main() {
    ...
    pthread_t p_threads[MAX_THREADS];
    pthread_attr_t attr;
    pthread_attr_init (&attr);
    for (i=0; i< num_threads; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void *) &hits[i]);
    }
    for (i=0; i< num_threads; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
    ...
}
```

```
#include <pthread.h>
#include <stdlib.h>
#define NUM_THREADS 32
void *compute_pi (void *);
...
int main(...) {
    ...
    pthread_t p_threads[NUM_THREADS];
    pthread_attr_t attr;
    pthread_attr_init(&attr);
    for (i=0; i< NUM_THREADS; i++) {
        hits[i] = i;
        pthread_create(&p_threads[i], &attr, compute_pi,
            (void*) &hits[i]);
    }
    for (i=0; i< NUM_THREADS; i++) {
        pthread_join(p_threads[i], NULL);
        total_hits += hits[i];
    }
}
```

default attributes

thread function

thread argument


# Thread Basics: Creation and Termination (Example)

```
void *compute_pi (void *s) {
    int seed, i, *hit_pointer;
    double rand_no_x, rand_no_y;
    int local_hits;
    hit_pointer = (int *) s;
    seed = *hit_pointer;
    local_hits = 0;
    for (i = 0; i < sample_points_per_thread; i++) {
        rand_no_x =(double) (rand_r(&seed))/(double) ((2<<14)-1);
        rand_no_y =(double) (rand_r(&seed))/(double) ((2<<14)-1);
        if (((rand_no_x - 0.5) * (rand_no_x - 0.5) +
            (rand_no_y - 0.5) * (rand_no_y - 0.5)) < 0.25)
            local_hits ++;
        seed *= i;
    }
    *hit_pointer = local_hits;
    pthread_exit(0);
}
```

# Programming and Performance Notes

- Note the use of the function `rand_r` (instead of superior random number generators such as `drand48`).
- Executing this on a 4-processor SGI Origin, we observe a 3.91 fold speedup at 32 threads. This corresponds to a parallel efficiency of 0.98!
- We can also modify the program slightly to observe the effect of false-sharing.
- The program can also be used to assess the secondary cache line size.

# Programming and Performance Notes



- Execution time of the `compute_pi` program.

# Synchronization Primitives in Pthreads

## Mutual Exclusion for Shared Variables

- Tasks work together to manipulate data and accomplish a given task.
- When multiple threads attempt to manipulate the same data item the results can often be **incoherent** if proper care is not taken to synchronize them
- Much of the effort associated with writing correct threaded programs is spent on **synchronizing concurrent threads with respect to their data accesses or scheduling**

The variable `my_cost` is thread-local and `best_cost` is a global variable shared by all threads.

```
1  /* each thread tries to update variable best_cost as follows */
2  if (my_cost < best_cost)
3      best_cost = my_cost;
```

- Assume that there are two threads,
- The initial value of `best_cost` is 100,
- The values of `my_cost` are 50 and 75 at threads `t1` and `t2`, respectively.
- If both threads execute the condition inside the `if` statement concurrently, then both threads enter the then part of the statement.
- Depending on which thread executes first, the value of `best_cost` at the end could be either 50 or 75.

## There are two problems here:

1. non-deterministic nature of the result;
2. more importantly, the value 75 of `best_cost` is inconsistent in the sense that no serialization of the two threads can possibly yield this result.

**Result of the computation depends on the race between competing threads**

- Critical segment- segment that must be executed by only one thread at any time.
- Threaded APIs provide support for implementing critical sections and atomic operations using

***mutex-locks***



# Mutual Exclusion

- **Critical segments** in Pthreads are implemented **using mutex locks**.
- Mutex-locks have two states:
  - locked and unlocked
  - At any point of time, only one thread can lock a mutex lock. A lock is an atomic operation.
- A thread entering a critical segment first tries to get a lock. It goes ahead when the lock is granted.

- **Mutex locks enforce mutual exclusion in Pthreads**
  - mutex lock states: locked and unlocked
  - only one thread can lock a mutex lock at any particular time
- **Using mutex locks**
  - request lock before executing critical section
  - enter critical section when lock granted
  - release lock when leaving critical section

created by  
`pthread_mutex_attr_init`  
specifies mutex type

# Mutual Exclusion

The Pthreads API provides the following functions for handling mutex-locks:

```
int pthread_mutex_lock (  
    pthread_mutex_t *mutex_lock);  
int pthread_mutex_unlock (  
    pthread_mutex_t *mutex_lock);  
int pthread_mutex_init (  
    pthread_mutex_t *mutex_lock,  
    const pthread_mutexattr_t *lock_attr);
```

# Mutual Exclusion

- We can now write our previously incorrect code segment as:

```
pthread_mutex_t minimum_value_lock;
...
main() {
    ....
    pthread_mutex_init(&minimum_value_lock, NULL);
    ....
}
void *find_min(void *list_ptr) {
    ....
    pthread_mutex_lock(&minimum_value_lock);
    if (my_min < minimum_value)
        minimum_value = my_min;
    /* and unlock the mutex */
    pthread_mutex_unlock(&minimum_value_lock);
}
```

function to initialize a mutex-lock to its unlocked state

- If the mutex-lock is already locked, the calling thread **blocks**
- otherwise the mutex-lock is locked and the calling thread **returns**

- a thread must unlock the mutex-lock
- If it does not do so, no other thread will be able to enter this section,
- (typically resulting deadlock.)

# Producer-Consumer Using Locks

## Constraints

- The producer-consumer scenario imposes the following constraints:
- The producer thread must not overwrite the shared buffer when the previous task has not been picked up by a consumer thread.
- The consumer threads must not pick up tasks until there is something present in the shared data structure.
- Individual consumer threads should pick up tasks one at a time.

```
pthread_mutex_t task_queue_lock;
int task_available;
...
main() {
    ...
    task_available = 0;
    pthread_mutex_init(&task_queue_lock, NULL);
    ...
}
void *producer(void *producer_thread_data) {
    ...
    while (!done()) {
        inserted = 0;
        create_task(&my_task);
        while (inserted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (work_available == 0) {
                consumer_work = my_task; work_available = 1;
                inserted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
    }
}
```

critical section



```
void *consumer(void *consumer_thread_data) {
    int extracted;
    struct task my_task;
    /* local data structure declarations */
    while (!done()) {
        extracted = 0;
        while (extracted == 0) {
            pthread_mutex_lock(&task_queue_lock);
            if (work_available == 1) {
                my_task = consumer_work;
                work_available = 0;
                extracted = 1;
            }
            pthread_mutex_unlock(&task_queue_lock);
        }
        process_task(my_task);
    }
}
```

critical section



# Mutex Types

```
pthread_mutex_init(&minimum_value_lock, NULL);
```

Mutex Types



- Normal
  - —thread deadlocks if tries to lock a mutex it already has locked
- Recursive
  - — single thread may lock a mutex as many times as it wants — increments a count on the number of locks —thread relinquishes lock when mutex count becomes zero
- Errorcheck
  - —report error when a thread tries to lock a mutex it already locked —report error if a thread unlocks a mutex locked by another



# Overheads of Locking

- Locks represent serialization points since critical sections must be executed by threads one after the other.
- Encapsulating large segments of the program within locks can lead to significant performance degradation.
- It is often possible to reduce the idling overhead associated with locks using an alternate function, `pthread_mutex_trylock`.

```
int pthread_mutex_trylock (pthread_mutex_t *mutex_lock) ;
```

- `pthread_mutex_trylock` is typically much faster than `pthread_mutex_lock` on typical systems
  - since it does not have to deal with queues associated with locks for multiple threads waiting on the lock.
  - enables a thread to do something else if a lock is unavailable

# Alleviating Locking Overhead (Example)

```
/* Finding k matches in a list */
void *find_entries(void *start_pointer) {
    /* This is the thread function */
    struct database_record *next_record;
    int count;
    current_pointer = start_pointer;
    do {
        next_record = find_next_entry(current_pointer);
        count = output_record(next_record);
    } while (count < requested_number_of_records);
}

int output_record(struct database_record *record_ptr) {
    int count;
    pthread_mutex_lock(&output_count_lock);
    output_count++;
    count = output_count;
    pthread_mutex_unlock(&output_count_lock);
    if (count <= requested_number_of_records)
        print_record(record_ptr);
    return (count);
}
```

# Alleviating Locking Overhead (Example)

```
/* rewritten output_record function */
int output_record(struct database_record *record_ptr) {
    int count;
    int lock_status;
    lock_status=pthread_mutex_trylock(&output_count_lock);
    if (lock_status == EBUSY) {
        insert_into_local_list(record_ptr);
        return(0);
    }
    else {
        count = output_count;
        output_count += number_on_local_list + 1;
        pthread_mutex_unlock(&output_count_lock);
        print_records(record_ptr, local_list,
            requested_number_of_records - count);
        return(count + number_on_local_list + 1);
    }
}
```

## ✓ Thread

- ✓ Basics
- ✓ The POSIX Thread API,
- ✓ Thread Creation and Termination,
- ✓ Synchronization Primitives in Pthreads,
- ✓ Controlling Thread and Synchronization Attributes,
- ✓ Thread Cancellation, Composite Synchronization Constructs.

## ✓ Thread

- ✓ Basics
- ✓ The POSIX Thread API,
- ✓ Thread Creation and Termination,
- ✓ Synchronization Primitives in Pthreads,
- ✓ Controlling Thread and Synchronization Attributes,
- ✓ Thread Cancellation, Composite Synchronization Constructs.

# Condition Variables for Synchronization

- A condition variable is a data object used for synchronizing threads
- This variable allows a thread to block itself until specified data reaches a predefined state
- Always use condition variables together with a mutex lock.

The shared variable `task_available` must become 1 before the consumer threads can be signaled.

The boolean condition `task_available == 1` is referred to as a predicate.

- The condition variables to atomically block threads until a particular condition is true.

# Condition Variables for Synchronization

- If the predicate is not true, the thread waits on the condition variable associated with the predicate using the function `pthread_cond_wait`.
- **`int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);`**
  - A call to this function blocks the execution of the thread until it receives a signal from another thread

**`int pthread_cond_signal(pthread_cond_t *cond);`**

1. Unblocks at least one thread that is currently waiting on the condition variable `cond`.
2. The producer then relinquishes its lock on `mutex` by explicitly calling `pthread_mutex_unlock`
3. allowing one of the blocked consumer threads to consume the task

**`int pthread_cond_broadcast(pthread_cond_t *cond);`**

wake all threads that are waiting on the condition variable

- Using a condition variable —
  - thread can block itself until a condition becomes true
    - thread locks a mutex
    - tests a predicate defined on a shared variable if predicate is false, then wait on the condition variable waiting on condition variable unlocks associated mutex
  - when some thread makes a predicate true
    - Thread can signal the condition variable to either wake one waiting thread wake all waiting threads
    - –when thread releases the mutex, it is passed to first waiter



# Condition Variables for Synchronization

**int pthread\_cond\_init (pthread\_cond\_t \*cond, const pthread\_condattr\_t \*attr);**



Initializes a condition variable (pointed to by cond) whose Attributes are defined in the attribute object attr

**int pthread\_cond\_destroy(pthread\_cond\_t \*cond);**



If at some point in a program a condition variable is no longer required,

**int pthread\_cond\_timedwait (pthread\_cond\_t \*cond, pthread\_mutex\_t \*mutex, const struct timespec \*abstime);**

1. Thread can perform a wait on a condition variable until a specified time expires.
2. At this point, the thread wakes up by itself if it does not receive a signal or a broadcast.

# Producer-Consumer Using Condition Variables

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;
/* other data structures here */

main() {
    /* declarations and initializations */
    task_available = 0;
    pthread_init();
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
}
```

# Producer-Consumer Using Condition Variables

```
void *producer(void *producer_thread_data) {  
    int inserted;  
    while (!done()) {  
        create_task();  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 1)  
            pthread_cond_wait(&cond_queue_empty, &task_queue_cond_lock);  
        insert_into_queue();  
        task_available = 1;  
        pthread_cond_signal(&cond_queue_full);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
    }  
}
```

# Producer-Consumer Using Condition Variables

```
void *consumer(void *consumer_thread_data) {  
    while (!done()) {  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 0)  
            pthread_cond_wait(&cond_queue_full, &task_queue_cond_lock);  
        my_task = extract_from_queue();  
        task_available = 0;  
        pthread_cond_signal(&cond_queue_empty);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
        process_task(my_task);  
    }  
}
```

# Controlling Thread and Synchronization Attributes

- The Pthreads API allows a programmer to change the default attributes of entities using *attributes objects*.
- An attributes object is a data-structure that describes entity (thread, mutex, condition variable) properties.
- Once these properties are set, the attributes object can be passed to the method initializing the entity.
- Enhances modularity, readability, and ease of modification.

`Int pthread_attr_init ( pthread_attr_t *attr);`

← initializes the attributes object attr to the default values

# Attributes Objects for Threads

- Use `pthread_attr_init` to create an attributes object.
- Individual properties associated with the attributes object can be changed using the following functions:

```
pthread_attr_setdetachstate,  
pthread_attr_setguardsize_np,  
pthread_attr_setstacksize,  
pthread_attr_setinheritsched,  
pthread_attr_setschedpolicy, and  
pthread_attr_setschedparam
```

# Thread Cancellation

```
int pthread_cancel ( pthread_t thread);
```

- When a call to this function is made, a cancellation is sent to the specified Thread
- The function returns a 0 on successful completion.

# Composite Synchronization Constructs

- By design, Pthreads provide support for a basic set of operations.
- Higher level constructs can be built using basic synchronization constructs.
- We discuss two such constructs - read-write locks and barriers.



# Read-Write Locks

- In many applications, a data structure is read frequently but written infrequently. For such applications, we should use read-write locks.
  - multiple reads can proceed without any coherence problems. However, writes must be serialized.
- A read lock is granted when there are other threads that may already have read locks.
- If there is a write lock on the data (or if there are queued write locks), the thread performs a condition wait.
- If there are multiple threads requesting a write lock, they must perform a condition wait.
- With this description, we can design functions for read locks `mylib_rwlock_rlock`, write locks `mylib_rwlock_wlock`, and unlocking `mylib_rwlock_unlock`.

# Read-Write Locks

- The lock data type `mylib_rwlock_t` holds the following:
  - a count of the number of readers,
  - the writer (a 0/1 integer specifying whether a writer is present),
  - a condition variable `readers_proceed` that is signaled when readers can proceed,
  - a condition variable `writer_proceed` that is signaled when one of the writers can proceed,
  - a count `pending_writers` of pending writers, and
  - a mutex `read_write_lock` associated with the shared data structure

# Read-Write Locks

```
typedef struct {
    int readers;
    int writer;
    pthread_cond_t readers_proceed;
    pthread_cond_t writer_proceed;
    int pending_writers;
    pthread_mutex_t read_write_lock;
} mylib_rwlock_t;

void mylib_rwlock_init (mylib_rwlock_t *l) {
    l -> readers = l -> writer = l -> pending_writers = 0;
    pthread_mutex_init(&(l -> read_write_lock), NULL);
    pthread_cond_init(&(l -> readers_proceed), NULL);
    pthread_cond_init(&(l -> writer_proceed), NULL);
}
```

# Read-Write Locks

```
void mylib_rwlock_rlock(mylib_rwlock_t *l) {
    /* if there is a write lock or pending writers, perform condition wait.. else
       increment count of readers and grant read lock */
    pthread_mutex_lock(&(l -> read_write_lock));
    while ((l -> pending_writers > 0) || (l -> writer > 0))
        pthread_cond_wait(&(l -> readers_proceed),
            &(l -> read_write_lock));
    l -> readers ++;
    pthread_mutex_unlock(&(l -> read_write_lock));
}
```

# Read-Write Locks

```
void mylib_rwlock_wlock(mylib_rwlock_t *l) {
    /* if there are readers or writers, increment pending writers count and
       wait. On being woken, decrement pending writers count and increment
       writer count */

    pthread_mutex_lock(&(l -> read_write_lock));
    while ((l -> writer > 0) || (l -> readers > 0)) {
        l -> pending_writers ++;
        pthread_cond_wait(&(l -> writer_proceed),
                          &(l -> read_write_lock));
    }
    l -> pending_writers --;
    l -> writer ++;
    pthread_mutex_unlock(&(l -> read_write_lock));
}
```

# Read-Write Locks

```
void mylib_rwlock_unlock(mylib_rwlock_t *l) {
    /* if there is a write lock then unlock, else if there are read locks,
       decrement count of read locks. If the count is 0 and there is a pending
       writer, let it through, else if there are pending readers, let them all go
       through */
    pthread_mutex_lock(&(l -> read_write_lock));
    if (l -> writer > 0)
        l -> writer = 0;
    else if (l -> readers > 0)
        l -> readers --;
    pthread_mutex_unlock(&(l -> read_write_lock));
    if ((l -> readers == 0) && (l -> pending_writers > 0))
        pthread_cond_signal(&(l -> writer_proceed));
    else if (l -> readers > 0)
        pthread_cond_broadcast(&(l -> readers_proceed));
}
```

# Barriers

- As in MPI, a barrier holds a thread until all threads participating in the barrier have reached it.
- Barriers can be implemented using a counter, a mutex and a condition variable.
- A single integer is used to keep track of the number of threads that have reached the barrier.
- If the count is less than the total number of threads, the threads execute a condition wait.
- The last thread entering (and setting the count to the number of threads) wakes up all the threads using a condition broadcast.

# Barriers


```
typedef struct {  
    pthread_mutex_t count_lock;  
    pthread_cond_t ok_to_proceed;  
    int count;  
} mylib_barrier_t;  
  
void mylib_init_barrier(mylib_barrier_t *b) {  
    b -> count = 0;  
    pthread_mutex_init(&(b -> count_lock), NULL);  
    pthread_cond_init(&(b -> ok_to_proceed), NULL);  
}
```



# Barriers

```
void mylib_barrier (mylib_barrier_t *b, int num_threads)
{
    pthread_mutex_lock(&(b -> count_lock));
    b -> count ++;
    if (b -> count == num_threads) {
        b -> count = 0;
        pthread_cond_broadcast(&(b -> ok_to_proceed));
    }
    else
        while (pthread_cond_wait(&(b -> ok_to_proceed),
                                &(b -> count_lock)) != 0);
    pthread_mutex_unlock(&(b -> count_lock));
}
```

# Barrier



- Execution time of 1000 sequential and logarithmic barriers as a function of number of threads on a 32 processor SGI Origin 2000.

# Implement DAXPY loop using pthread(32threads)

## 4. DAXPY Loop:

The daxpy loop is the core of the benchmark. This loop is used to measure the performance. By using this loop we can observe how fast a certain machine can execute. Daxpy loop multiplies a vector by a scalar and adds it to another vector.

D stands for Double precision, A is a scalar value, X and Y are one-dimensional vectors of size 216 each, P stands for Plus. The operation to be completed in one iteration i.e  $X[i] = a * X[i] + Y[i]$ .

```
void daxpy(double y[], double a, double x[], int n)
{
    int i;
    for (i = 0; i < n; i++)
        y[i] = a*x[i] + y[i];
}
```